

The Development of a Collaborative Tool to Teach Debugging

Sammy Furr

December 3, 2020

Abstract

TODO: write an abstract

Contents

1	Introduction	4
1.1	Motivation	4
1.1.1	The Value of Teaching Debugging	5
1.1.2	Methods for Teaching Debugging	6
1.1.3	The Value of Collaborative Programming	7
2	Requirments	9
2.1	Encouraging Collaboration	9
2.2	Enhancing Student's Understanding	10
2.3	Existing Tools that Enable Collaborative Programming	10
3	Tools Used	10
3.1	Kubernetes	12
3.2	Mozilla's rr	12
3.2.1	Overview	12
3.2.2	Limitations	14
3.3	pygdbmi	14
3.4	Socket.IO	15
3.5	MongoDB	15
3.6	Flask and Node.js	16
3.7	React	16

3.8	Monaco and Xterm.js	17
4	Design	18
4.1	Configuration and Setup	20
4.1.1	Cluster Selection and Configuration	21
4.1.2	Creating Pods	21
4.1.3	Service Creation	25
4.2	RR Debug Sessions & RR Message Server	26
4.2.1	RR Message Server	28
4.2.2	RR Debug Session	29
4.3	Frontend/API Server	35
4.3.1	Frontend/API Server Configuration	36
4.3.2	The API Server	37
4.4	Webapp	45
5	Next Steps	46

1 Introduction

The goal of this project is to create a collaborative debugger to aid in the teaching of debugging. The collaborative debugger aims to help teachers integrate this undertaught skill into their classes. It realizes the benefits of collaborative programming and teaching debugging by providing a platform that is genuinely useful to students and teachers.

1.1 Motivation

Debugging is invaluable in writing and understanding code, yet it is rarely formally taught [20]. We typically teach students programming structures, concepts, and languages, but leave them to learn the tools they use to write code by themselves. This approach often works well—a programmer’s choice of tools is often *very* personal and students figure out how to configure an individualized workflow. Perhaps because debuggers are tools, students are often expected to learn them with minimal guidance. Unlike editors or reference guides however, effectively using a debugger requires a set of high-level, platform agnostic, teachable skills. Teaching these skills is effective, and translates into better, faster, debugging and programming [19] [22]. Teaching these debugging skills collaboratively will likely offer the same confidence and correctness benefits realized by teaching programming collaboratively [18] [21].

1.1.1 The Value of Teaching Debugging

There is an unfortunate lack of research specifically into the efficacy of teaching debugging for computer science students, despite a recent rise in the inclusion of debugging in “computational thinking” curriculums [22]. These curriculums attempt to teach skills in computer science classes that translate into other subject areas: the UK’s computer science curriculum considers debugging an essential “transferable skill” [13].

There seems to be confidence that the problem-solving techniques used in debugging are widely applicable, but of greater interest to computer science teachers is whether teaching debugging directly benefits student programmers. Michaeli and Romeike conducted a good, albeit somewhat small, study on the efficacy of teaching a systematic debugging process to K12 students. They found that students who have been taught a specific debugging framework performed better in debugging tests and were more confident in their own debugging skills [22]. Their result is positive evidence towards the efficacy of teaching debugging, though it doesn’t include college or university students.

As Michaeli and Romeike point out, there is a lack of research into the value of teaching debugging in higher education. None of the research these authors found placed much focus on explicitly teaching debugging. Chmiel and Loui studied whether students who were provided with debugging tools and frameworks performed better on tests or spent less time on assignments

than those who were not [14]. Though this research wasn't able to find conclusive evidence towards better performance on tests or assignments, it did find that students in the treatment group felt more confident in their debugging abilities. Unfortunately Chmiel and Loui's study didn't involve extended explicit teaching of debugging—use of the tools was voluntary, and variations in the students' individual abilities made the data difficult to evaluate.

Though there is a lack of higher-education research, the value of teaching debugging is still demonstrable. The research discussed all finds that K-12 and college students alike commonly resort to sporadic debugging techniques when beginning to learn. Since this pattern of behavior that explicitly teaching debugging corrects exists in college as well as in K-12 students, it seems logical that the benefit of explicitly teaching debugging to K-12 students should be realized equally by their collegiate counterparts.

1.1.2 Methods for Teaching Debugging

Similarly to research on the value of teaching debugging, research into how to best teach debugging is sparse. Chan et al. allow that “in general research on how to improve debugging is sporadic”—an observation that leads them to research a framework to reduce the complexity of teaching debugging [19]. To organize their framework, they split debugging knowledge into 5 categories: *Domain*, *System*, *Procedural*, *Strategic*, and *Experiential*. They then review different debugging tools and teaching aids—from those that involve writing

code to games—and map tools to the knowledge areas they seek to address. After an evaluation of a host of different tools, they claim to find a few significant faults in current debugging teaching platforms. The collaborative debugger primarily seeks to address the lack of back-tracing ability/coverage found in their research.

1.1.3 The Value of Collaborative Programming

Research into the value of collaborative programming is overwhelmingly positive. McDowell et. al. found that not only does collaborative programming significantly boost student confidence and the retention of students in computer science majors, but that it demonstrably improves student’s work [21]. These benefits of confidence and correctness are present when paired and non-paired students are given identical assignments [18], further indicating that the simple act of collaboration definitively benefits computer science students.

It seems that the benefits in confidence that result from teaching debugging should be magnified by teaching debugging collaboratively. There are similarities between the introductory nature of teaching fundamental debugging skills and the nature of teaching fundamental programming skills covered in the classes studied in [18] and [21]. Introductory programming classes typically use a specific programming language in order to introduce widely applicable programming principles. By using a specific collaborative platform to introduce debugging principles, students may realize the same

benefits in both confidence *and* correctness that they do from pair programming in introductory computer science classes.

Debuggers exist at an intersection of tools and skills similar to programming languages themselves. By becoming familiar with a specific debugger, students may learn techniques and paradigms necessary to use all debuggers effectively. The collaborative debugger aims to provide the optimal platform for students to learn debugging skills.

2 Requirments

In order to provide the best platform for both teachers and students, the collaborative debugger must fupill two key requirements:

1. It must encourage collaboration through a seamless experience.
2. It must make it easier to enhance students' understanding of debugging and programs being debugged.

2.1 Encouraging Collaboration

A large number of tools exist to facilitate collaborative programming. COVID19 has greatly increased the demand for tools that not only make collaboration easier, but make it easy remotely. The tool that was most influential in the design of the collaborative debugger, Replit [25], enables remote collaborative programming.

Replit provides a simple browser-based IDE for over 50 languages. Programmers choose a language and can edit and run code collaboratively inside the Replit webapp. Similarly to other collaborative text-editors, input and output from all users is synced.

The collaborative debugger aims to provide a similar experience to Replit. Students should be able to start a debugging session, easily invite other students, and interact with the debugger in a way that makes it easy to share their knowledge. By allowing multiple users to interact with the same debug-

ging session together in real-time, the collaborative debugger lets students share experience and work through problems together. The debugger also makes it easy for teachers to demonstrate debugging techniques remotely to multiple students, any of whom can also interact with the debugger. Since debugging sessions are hosted remotely, users can start a session in one location and resume it later from a different computer.

2.2 Enhancing Student's Understanding

Apart from the benefits to understanding that are realized through collaboration, the collaborative debugger should enhance students' understanding of debugging in ways that a traditional debugger cannot.

The aim is to provide students with an experience they can directly reference when using `rr` or `gdb` in the future, as well as one which is seamless to encourage the acquisition of debugging skills.

2.3 Existing Tools that Enable Collaborative Programming

3 Tools Used

The next sections give an overview of the various tools used to create the collaborative debugger. The debugger consists of:

1. A frontend web app built using React (3.7) that presents a debugging

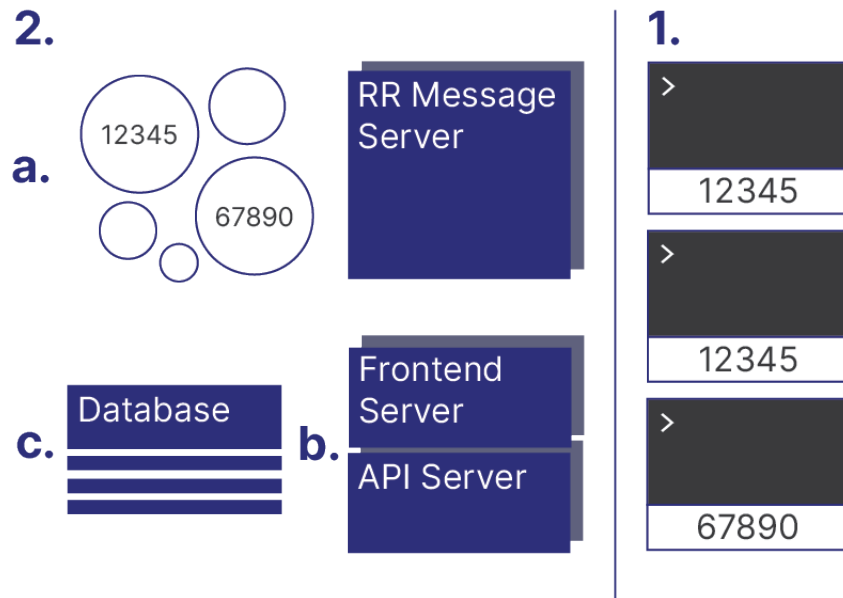


Figure 1: Overview of the Collaborative Debugger

interface to the end user.

2. A distributed backend managed by Kubernetes (3.1), split into three parts:
 - (a) A Pod for each debugging instance which runs the rr debugger (3.2). These communicate directly with users through WebSocket server Pods. (3.4).
 - (b) Pods running a frontend server written in Node.js which works in tandem with an API server created using Flask (3.6). The API server manages creation and destruction of debug sessions, as well as authentication.
 - (c) A MongoDB (3.5) database.

3.1 Kubernetes

Kubernetes is the defacto standard in container orchestration software. It provides a layer of abstraction on top of normal containers, like those created by Docker. By bundling one or more closely linked containers into a “Pod”, Kubernetes is able to manage deployment and re-deployment of applications running inside containers. It is trivial to create new Pods, or to create multiple Pods running the same application as needed within a Kubernetes cluster [7]. The speed at which even relatively large Pods can be created and the inherent security provided by containerization drove the decision to create a new Pod on the fly for each debugging instance in the collaborative debugger. This allows the debugger to provide a similar level of convenience to existing collaborative tools, such as repl.it (2.3).

Kubernetes also provides services to facilitate load balancing, manage storage volumes, and contain secrets. The abstraction provided by these features, in tandem with the ease of Kubernetes deployment on a managed Kubernetes service [2] greatly accelerated development.

3.2 Mozilla’s rr

3.2.1 Overview

rr is “a lightweight tool for recording, replaying and debugging execution of applications” [24]. rr allows a programmer to record the execution of a program on any compatible machine and replay the execution later. This

enhances GDB’s ability to “time-travel” when debugging, using commands such as `reverse-continue` and `reverse-stepi` [16] to step backwards and forwards through a program’s execution. Through a novel encapsulation of the execution space, `rr` is able to deterministically record and replay the execution of syscalls and other process behavior that differs run-to-run. This is invaluable when trying to debug behavior that is not entirely dependent on the code being debugged. A typical workflow in `rr` consists of recording an inexplicable error, replaying execution to find the area in which the error occurs, and then narrowing in on the bug not by re-running the entire program, but by progressing back and forth through execution in the problem area.

`rr` is an ideal tool for teaching debugging because it allows instructors to record execution of a program and design a debugging example with the knowledge that normally non-deterministic events will be repeatable, and that any input they provide to the program will be exactly replicated. With the collaborative debugger, teachers can record a program’s execution and design a debugging lesson which students can work on together. The repeatability of `rr` means that students can focus on debugging, and teachers can create as specific examples as they please. The use of `rr` is the most significant step the collaborative debugger takes to addressing the lack of back-tracing ability/coverage found in existing tools for teaching debugging [19].

3.2.2 Limitations

In comparison to solutions like PANDA [15] that rely on capturing the entire state of a virtual machine to replay execution, rr records and replays faster, produces far smaller files, and doesn't force execution inside of a VM. [23] The trade off for these benefits are two major system limitations: rr is only compatible with the Linux kernel, and it's deterministic recording and replay relies on a feature that is only found on modern *Intel* x86 CPUs. These limitations influenced the development of this project as a webapp similar to existing tools for collaborative programming.

Luckily, the speed and size benefits of rr lend themselves well to non-local execution. In conjunction with Kubernetes, it takes a few seconds to create a new container running rr and connect to web clients.

3.3 pygdbmi

In order to “support the development of systems which use the debugger as just one small component of a larger system”, GDB provides a machine-oriented interface called GDB/MI [16]. rr supports interaction through GDB/MI, and using the interface was a natural choice for the collaborative debugger. In addition to being far easier to interact with from within a program, the structured, machine-friendly output of GDB/MI lends itself in particular to future development of visualization aids in the collaborative debugger.

To parse rr output into Python dictionaries and to easily control rr as a

subprocess, pygdbmi [26] is used in each debugging Pod. pygdbmi’s abstraction simplifies programatically controlling rr. A Pod can receive a command from the client, pass it to rr, and respond without having to deal with parsing GDB/MI output or with directly managing the rr process.

3.4 Socket.IO

To speed communication, the collaborative debugger uses WebSockets to directly connect web clients and the Pods running rr. Socket.IO is a library that extends WebSockets. It provides backup incase a WebSocket connection cannot be established, enables automatic reconnection and disconnection detection, and adds support for namespaces [11]. The collaborative debugger uses the standard JavaScript implementation of Socket.IO on the client side. Messages are passed through a server to individual debugging Pods, both of which use the Python implementation of Socket.IO, python-socket.io [17].

3.5 MongoDB

The collaborative debugger uses a database to store information about users, Pods, and example debugging sessions. Due to it’s speed of deployment and natural interaction with the object-oriented languages used to create the project, MongoDB was chosen as database software [3].

3.6 Flask and Node.js

The primary server for the collaborative debugger is split into two sections: a simple Node.js [9] server that serves the frontend webapp, and an API server created using Flask [6]. While in development, the builtin React (3.7) development server is used to serve the frontend. This makes debugging the frontend far easier.

An API server is necessary to authenticate users and to provide a means to create/delete debugging sessions. Since the rest of the backend was created using Python, Flask was chosen to create the API server. Flask is a lightweight web application framework which lends itself perfectly to interacting with the Python MongoDB and Kubernetes APIs.

3.7 React

React is JavaScript library that simplifies creating user interfaces and managing state [10]. React's state management is of particular importance to the collaborative debugger's frontend. State constantly changes as users create/delete debugging sessions, join existing sessions, and communicate with rr. React allows classes to encapsulate components such as a list of existing debugging sessions, a view of the current program's source code, and the terminal interface with rr. Instances of these classes maintain state and update efficiently.

The frontend makes extensive use of JSX, syntax which allows the inclu-

sion of segments of HTML code within a React app written in JavaScript. This makes it easy for each component of the one-page webapp to hide/show subcomponents as state changes.

3.8 Monaco and Xterm.js

After joining or creating a debugging session, users spend most of their time interacting collaboratively with rr. Their primary interface to rr is through Xterm.js, a frontend component that makes it easy to emulate terminal behavior in the browser [12]. With a few control methods, it is simple to provide a terminal interface to rr that is virtually indistinguishable from a local session. By using the Xterm.js based interface, students can learn to use rr (and by extension gdb) collaboratively, and directly translate that knowledge to individual work.

In addition to the terminal interface, the frontend shows a view of the current source file being debugged. The Monaco Editor [8] is used to display this source view. Though more complex than is strictly necessary to display code, Monaco makes it easy to format and syntax-highlight. Using Monaco also simplifies the future addition of editing source code, should the need arise. React’s state management allows updating text in the editor as efficiently as possible.

4 Design

The collaborative debugger consists of a distributed Kubernetes backend and frontend React webapp. Kubernetes was chosen for the backend primarily so that a Pod could be created dynamically for each debugging session. The design of the backend is heavily distributed, allowing individual components to be modified without the whole system needing to be reconfigured.

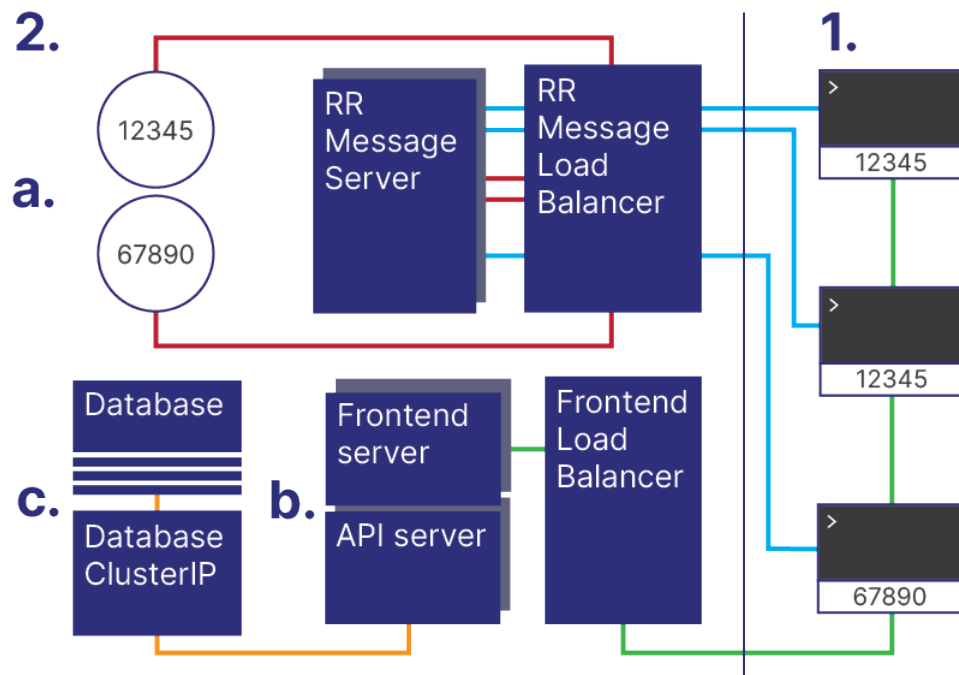


Figure 2: Detailed Overview of the Collaborative Debugger

Each backend component of the collaborative debugger runs in its own individual Pod. There are two different classes of Pods in the collaborative debugger:

1. Statically created Pods. These are the *RR Message Server Pods*, the *Frontend/API Server Pods*, and the *Database Pod*. This class of Pods are manually created when the cluster that will run the backend is first initialized. The *RR Message Server Pods* and *Frontend/API Server Pods* may be created using Deployments [1] to allow later scaling, where multiple Pods running the same application may be created to facilitate increased load.
2. Dynamically created Pods. This class of Pod contains the individual instances of *RR Debug Session Pods* that are created on request by the API server. When a client requests a new debug session, the API server uses the Kubernetes API to create a new Pod based on an existing template, gives the Pod a unique identifier, and associates it in the database with the requesting client.

These Pods communicate with other Pods in the cluster and with the outside world through Services. The Kubernetes documentation defines a Service as “an abstraction which defines a logical set of Pods and a policy by which to access them” [1]. In the collaborative debugger, these Services manifest as:

1. The *Database ClusterIP*: a ClusterIP, which exposes the Database Pod only inside the cluster. The only component that makes use of this ClusterIP is the API server, which uses it primarily to communicate information about users and *RR Debug Sessions* with the database.

2. The *RR Message Load Balancer*: a Load Balancer, which exposes the RR Message Server Pods to the outside world. Using Socket.IO, clients send commands to and receive responses from individual RR Debug Session Pods through the RR Message Server Pods.
3. The *Frontend Load Balancer*: another Load Balancer, which exposes the Frontend/API Server Pods to the outside world. Clients request the frontend webapp and send api requests/receive api responses through this Load Balancer.

The frontend webapp dynamically updates as the user requests new debug sessions, issues commands to rr, and visits new source files. A user can be part of multiple debug sessions simultaneously. Each debug session is assigned a 5 digit identifier at creation, which is used to join sessions in progress.

Each component of the backend and frontend will be discussed in depth in the following sections.

4.1 Configuration and Setup

Configuration and setup of the collaborative debugger is relatively simple. After a Kubernetes cluster is created (this is made easier by using a Managed Kubernetes service) Pods and Services are created using various configuration files. Services should be created first, so that Pods which rely on access to Services function properly on creation. The following sections outline the process of creating a cluster, Pods, and Services, with a focus on statically

created Pods. The process of building images for building dynamically created Pods is similar, but the process of starting the Pods is more complicated. This process will be discussed in-depth in the section on the API server (4.3).

4.1.1 Cluster Selection and Configuration

Though Kubernetes aims to be largely platform-agnostic, the requirements of rr impose some restrictions on cluster setup and configuration. Clusters, even those running inside a VM, must be run on machines using relatively modern Intel x86 CPUs (Nehalem and beyond). The clusters must run on an operating system using Linux kernel version 3.11 or higher [24]. Finally, in order for rr to be able to work efficiently, the `kernel.perf_event_paranoid` parameter must be set to 1 [24]. This should be done on every node in the cluster which will run RR Debug Sessions. For the purposes of development, it has been set to 1 on all nodes in the collaborative debugger cluster.

4.1.2 Creating Pods

Pods are created in five steps:

1. A `Dockerfile` is used to build a new Docker container image from various pieces of source code, scripts, and a base image (such as the official MongoDB image or official Ubuntu image). The Dockerfile also contains instructions to install necessary packages, run build scripts, and create file structures in the image.




Figure 3: The Pod Creation Process

2. The Docker image is tagged and uploaded to a private container registry.
 3. A Pod configuration schema is defined/updated with details of the corresponding image's tag and any necessary Pod-specific settings/startup commands.
 4. The Pod configuration schema is applied, either statically or dynamically. When the schema is applied, Kubernetes pulls the image from the container registry and creates a new Pod according to the schema, running any startup commands if provided.
 5. If Pod creation is successful, the result is a new Pod running in the cluster.
- 1. Docker Container Creation** Docker containers are created using a Dockerfile. The Dockerfile used to create the container image for the RR

Message Server is shown below:

```
# Base Image
FROM ubuntu:latest

# Package Installation
WORKDIR /tmp/
ENV DEBIAN_FRONTEND="noninteractive"
RUN apt-get update && apt-get install -y \
python3-pexpect python3-pip

# User Creation
RUN useradd -ms /bin/bash rrserver
USER rrserver

# File structure creation/app setup
RUN pip3 install requests python-socketio \
eventlet
WORKDIR /home/rrserver/
RUN mkdir app
WORKDIR /home/rrserver/app/
COPY server.py .
COPY startup.sh .

# Startup command
CMD ["sh", "startup.sh"]
```

Listing 1: RR Message Server Dockerfile

The build process for each collaborative debugger Docker image follows the same structure as the one outlined in the Dockerfile above:

1. The base image is defined. The RR Message Server and RR Debug Session images are based on the latest Ubuntu image. This is particularly necessary for the RR Debug Session image, as rr's low-level nature necessitates somewhat frequent updates as changes are made to the Linux kernel. The Frontend/API Server image is based on the latest Node image, and the Database image is based on the latest MongoDB image.
2. Second, any necessary packages are installed. For the RR Debug Pod image, rr is compiled from source and installed.

3. A non-root user is created if necessary.
4. Program files are copied over and a file structure is created. Packages that don't rely on the base images builtin package manager are installed at this time. In the example above, these include Python packages.
5. A startup command is defined.

Each line in a Dockerfile corresponds to a layer in the built image. This build order minimizes the amount of rebuilding necessary by placing the items that are most likely to change towards the end of the build process.

2. Container Registry Upload Most Managed Kubernetes services come with the option to create a private container registry. With proper authentication, this allows Docker and Kubernetes to access user-created images as easily as if they were in a public registry. Images built with Docker are uploaded to a private container registry for use in the collaborative debugger.

3. Pod Configuration Schema The Pod configuration schema for most Pods in the collaborative debugger is fairly generic. It consists of a **name**, an **image** sourced from the container registry, and in the case of pods that need to interact with a load balancer, an **app**.

```
apiVersion: v1
kind: Pod
metadata:
  name: rr-translation
  labels:
    purpose: translate-rr
```

```
spec:
  containers:
  - name: rr-test-container
    image: example-container-registry.com/sproj/...
    securityContext:
      capabilities:
        add:
        - SYS_PTRACE
    restartPolicy: OnFailure
```

Listing 2: RR Debug Session Schema

A notable exception is the RR Debug Session Schema, which adds the `SYS_PTRACE` capability to the Pod. This is necessary for `rr` to properly trace system calls.

4 & 5. Pod Creation For statically created Pods, the `kubectl apply` command is used to create new Pods. Kubernetes pulls the container image defined in the schema from the container registry and starts the Pod with any necessary commands. The database Pod is connected to a long-term storage volume at this time. Upon successful creation, the Pod is ready to interact with any necessary load balancers.

4.1.3 Service Creation

Services are the first components of the collaborative debugger created after the cluster is initialized. The three Services used by the debugger are all statically created. Like statically created Pods, schemas that define Services are applied manually using the `kubectl apply` command.

```
apiVersion: v1
kind: Service
metadata:
```

```
name: frontend-load-balancer
spec:
  selector:
    app: frontend
  type: LoadBalancer
  ports:
    - port: 3000
      targetPort: 3000
```

Listing 3: Frontend Load Balancer Schema

The `app` field in the above schema corresponds to the `app` field defined in the metadata of the Frontend/API Server configuration schema. Traffic to the Load Balancer’s external IP address on port 3000 is redirected automatically by Kubernetes to any Pod running the Frontend/API Server. The Load Balancer determines which Pod is the most suitable given current demands on the system.

4.2 RR Debug Sessions & RR Message Server

At the heart of the collaborative debugger is the RR Debug Session. Every time a user wants to debug a new program execution, an RR Debug Session Pod is dynamically created by the API server. The new Pod is assigned a unique five digit identification number when it is created. This five digit number, is used as the channel for the RR Debug Session, separating its communication from other RR Debug Session Pods. Clients, RR Message Server Pods, and the RR Debug Session Pod connect through the RR Message Load Balancer. Clients and RR Debug Sessions send messages using the channel assigned to the pod at creation time. A diagram of one full communication cycle between two clients and an RR Debug Session Pod is outlined in the

diagram below:

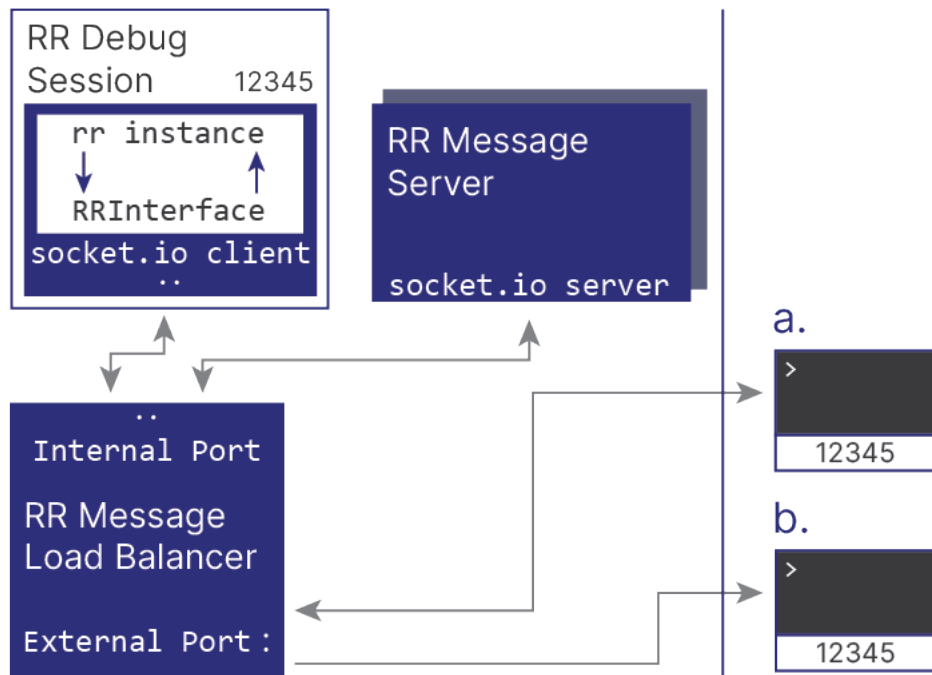


Figure 4: RR Message Server Communication Cycle

The process of sending a command to an RR Debug Session and receiving a response is as follows:

1. Client A, already having connected to the channel that corresponds to its current RR Debug Session (channel 12345 in the example), sends an `rr` command.
2. The RR Message Server receives the command and emits a message that RR Debug Session Pods are equipped to receive on the same channel. In practice, since only one RR Debug Session Pod is ever on a channel, this equates to emitting a message to the Pod directly.

3. The RR Debug Session Pod receives the message, and passes the command to its instance of RRInterface. When it receives a response, it emits the output from rr along with other debugging information and the channel.
4. The RR Message Server receives the response, and emits a response message that clients are equipped to receive on the corresponding channel.
5. Clients A and B are connected to the corresponding channel, so they both receive the response.

4.2.1 RR Message Server

The RR Message server is remarkably simple, consisting of just 3 important functions:

```
@sio.on('join_channel')
def join_channel(sid, data):
    sio.enter_room(sid, data['channel'])

@sio.on('rr_command')
def on_rr_command(sid, data):
    data['sid'] = sid
    try:
        sio.emit('rr_command', data,
                  room=data['channel'])
    except:
        pass

@sio.on('rr_response')
def on_rr_response(sid, data):
    sio.emit('rr_response', data,
             room=data['channel'])
```

Listing 4: RR Message Server

Even in a fully production ready version of the collaborative debugger with more security features enabled, the RR Message Server is unlikely to become much more complex. It exists purely to pass messages between clients and RR Debug Session Pods and to manage channels. Since the design of other aspects of the collaborative debugger ensure that all members of a given channel exist only during the lifetime of it's corresponding RR Debug Session Pod, the `'join_channel'` event handler simply adds socket.io clients to a specified channel on request.

The two message processing functions, `on_rr_command` and `on_rr_response`, are equally simple. When the RR Message Server receives an `'rr_command'` message, it passes the message data along to the corresponding RR Debug Session Pod by emitting an `'rr_command'` namespaced to the correct channel with `room=data['channel']`. Since the socket.io client running in RR Message Server Pods has an event handler defined for `'rr_command'`'s but frontend clients do not, only RR Message Server Pods receive the command. The same process happens in reverse for `'rr_response'`'s, with only the frontend socket.io clients having a handler defined for `'rr_response'`.

4.2.2 RR Debug Session

Building Example Pods The process for building RR Debug Session example images is differs slightly from other images used by the collaborative debugger. Currently, three example RR Debug Session images are available for use. Example container images are based on the primary RR

Debug Session image, RR Translation. The build process for this image, based on the latest official Ubuntu container image, installs rr, python, and all necessary packages as well as the app that will run on the final image, `rrtranslation.py`.

`FROM` example-container-registry.com/sproj/...

`WORKDIR` /home/debug/app/

`COPY` hash.c .

`COPY` names .

`COPY` startup.sh .

Listing 5: RR Debug Session Hash Example—Dockerfile

The Dockerfile shown above is for a simple hash program RR Debug Session example. The build process copies over any necessary files, as well as a startup script:

```
gcc -g -o hash hash.c
rr record ./hash
python3 rrtranslation.py $1
```

Listing 6: Example Startup Script

This script is used to start the `'rrtranslation.py'` program when the Pod is dynamically created. The script will be passed the Pod's channel as argument `$1` by the Kubernetes API on startup.

Current example pod startup scripts compile the program to be debugged

and record execution when the pod is created. It is easy to create a more specific example to be replayed, if exact reproduction of syscalls or other non-deterministic behavior is desired. The example's creator can record an rr session locally ahead of building the example Pod and copy the recording into the appropriate folder during the build process. Since rr recordings are portable, the example will replay without issue in the RR Debug Session Pod.

The RR Translation Program The RR Translation program consists of two main components: a socket.io client, `sio`, and an instance of the `RRInterface` class, `rr_i`, that controls the rr subprocess and parses interactions. Simple code to read information about the current source file currently exists within `sio`'s `'rr_command'` event handler, but should be broken out into its own class if more complexity is added.

The RR Translation program's socket.io client instance begins by connecting to the RR Message Server. It immediately emits a `'join_channel'` message, using the channel passed in from the Kubernetes API call on its creation.

```
if __name__ == '__main__':
    channel = sys.argv[1]
    sio.connect('http://rr-message-server-load-balancer:8000')
    sio.emit('join_channel', {'channel': channel})
```

Listing 7: RR Translation Main

After joining the appropriate channel, `sio` waits to receive an `'rr_command'`.

The first part of `sio`'s `'rr_command'` event handler is shown below:

```
@sio.on('rr_command')
def on_rr_command(data):
    body = {'from': data['sid'],
            'command': data['command'],
            'channel': channel}

    try:
        response = {'output': rri.write(data['command'])}
```

Listing 8: RR Command Event Handler

The handler first builds part of the response body, passing back data about the channel and originator of the message that will be important for the RR Message Server and client later. Next, it calls the function necessary to pass a message to rr and receive a response, `rri.write()`.

When the program starts, an instance of the `RRInterface` class, `rri`, is initialized in the same scope as `sio`. `RRInterface` contains an instance of the `GDBController` class from `pygdbmi` to interface with rr and control the rr subprocess, as well as a collection of methods to parse rr output, the three most important of which are shown below:

```
def get_full_rr_response(self, command):
    response = self.gdbmi.write(command)
    while(not(self.end(response) or self.exited(response))):
        response.extend(self.get_rr_response())
    return response

def write(self, command):
    if self.command_forbidden(command):
        raise DissallowedError
    self.timeline.append(command)
    return self.console_output(
        self.get_full_rr_response(command))

def source(self):
    f = self.get_full_rr_response(
```

```
        '-file-list-exec-source-file')[0]['payload']  
    return {'file': f['file'],  
            'line': f['line'],  
            'path': f['fullname']}
```

Listing 9: RRInterface

When `write` is called, it determines if the command to be executed is forbidden. Currently only one gdb command, `shell`, is disallowed. Parsing the output from `shell` is difficult, it is of virtually no use since *recordings*, not currently executing programs that `shell` could effect are being debugged, and it opens the door to security issues. Though containerization means that the most users could probably do is render their own debug session useless, the downsides of `shell` far outweigh the benefits.

Next, `write` appends the command to the `RRInterface`'s `timeline` instance variable. `timeline` is a list of all commands the debug session has executed. Though unused at the moment, it can serve in the future to implement a shared history between all clients and to facilitate the saving of debug sessions in-progress. By issuing all commands in `timeline` to a new instance of `RRInterface` with the same recording, it is possible to restore a debug session to an identical previous state. To save a debug session, the recording and `timeline` can be stored in a database, and be used to initialize a new Pod with identical state to a previous RR Debug Session Pod. This appears to clients as a seamless restoration of a previous Debug Session.

Since the frontend currently consists of a terminal interface to `rr` and a view of the current source code, `write` does not need to return any information from `rr` other than console output. Future versions of the col-

laborative debugger that support visualization tools could use gdbmi commands such as `-stack-list-frames` to retrieve information to be used by frontend visualization tools. `console_output` extracts user-relevant output from the dictionary returned by `GDBController`'s `write` method. The `get_full_rr_response` method ensures that all relevant output has been received from the rr subprocess before returning. Write reduces the lengthy dictionary that would be produced by a command as simple as `break main` into a few lines of user-relevant output.

After the `'rr_command'` event handler has received a response from `rr`, it calls `RRInterface`'s `source` function to retrieve information about the current source file being debugged. Most often the source file has not changed, and the only relevant piece of information that needs to be passed back to the client is the new line number in the source file. If the source file has changed, the handler reads its contents, updates `current_file`, and emits the line number, file name, and contents. If an error occurs at any point in the process, a detailed trace is emitted for debugging purposes. The trace should be omitted in production.

Most of the complexity in the RR Translation program stems from parsing rr output. The flow of data in the program is quite simple: a command is received, it is passed to rr, rr's response is processed, and a response message is emitted. This flow would remain unchanged even with the addition of functions to retrieve information for data visualization. Since the process of turning rr responses into a terminal interface takes place in the frontend,

rather than the RR Translation program providing some sort of REPL over WebSockets, updates can be made to the frontend webapp without requiring backend changes. The whole process is also quite snappy, with only a slight delay from the client's perspective compared to running a debugger natively.

Effort is taken in the design of the RR Debug Session and RR Message Server to ensure students receive a pleasant and near-identical experience to debugging locally in the future. The system of inheritance used to create example debug session pods is designed to be as simple as possible, though a more fully featured frontend could simplify this process further in the future.

4.3 Frontend/API Server

The second major backend system consists of the Frontend/API Server and database. This system is responsible for serving the frontend webapp, as well as processing all API requests, such as those to authenticate users and create new RR Debug Sessions. The builtin React development server is used to serve the frontend webapp and redirect API requests to the Flask API server during development. For both security and performance reasons this should be changed to a combination of a more robust solution like Nginx [4] and a production suitable web server for production. The relationships between components and overall process will remain unchanged after this migration.

Since the entire webapp is one page, the process for serving it is completely standard. The server receives a request for the index page of the website, and returns the compiled React webapp that makes up the homepage. The

process for API requests is more involved, and an example is outlined below:




Figure 5: API Request Process—Login

1. The client makes a POST request to the `'/login'` URL. This request contains the necessary data to log in the user.
2. The frontend server forwards the request of a URL it does not recognize to the API server running in the same Pod on a different port.
3. The API server recognizes the `'/login'` route, and communicates with the database to attempt to log in the user.
4. The API server returns the result of the login request, which is returned by the frontend server to the client.

4.3.1 Frontend/API Server Configuration

The Frontend/API Server uses Yarn [5] to manage packages. When a Frontend/API Server Pod is deployed, the Pod's startup script uses the `yarn start`

and `yarn start-api` commands to start the frontend server and API server respectively. These scripts are defined in the server's `package.json` configuration file:

```
"start": "react-scripts start",  
"start-api": "cd api && flask run --no-debugger",  
.  
.  
"proxy": "http://localhost:5000",
```

Listing 10: Frontend/API Server Configuration

Also of note in `package.json` is the instruction to proxy the API server, which is running on port 5000. This achieves the automatic forwarding of API requests to the API server described above. Since API requests are proxied through the same address serving the frontend webapp, there are no issues with cross-origin requests.

4.3.2 The API Server

The API server is implemented using Flask. The server program consists of handlers for the various API routes and instances of two classes which communicate with the database/Kubernetes API: `TranslationPodManager` and `UserManager`. The full structure of the API Server program, `api.py` is shown below:

Care was taken to abstract as much as possible in the design of the API server and the classes it instantiates. The two classes instantiated by `api.py`, `TranslationPodManager` and `UserManager`, each instantiate another class which interacts with the database, `DatabaseController`, as an instance vari-

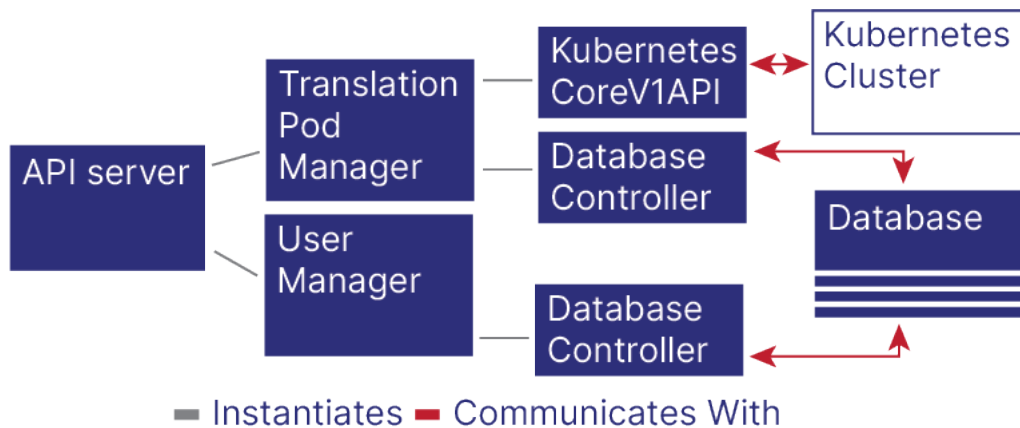


Figure 6: Structure of the API Server

able. `TranslationPodManager` further instantiates the Kubernetes Core V1 Python API to communicate with the cluster in order to create/delete pods. This process of abstraction ensures that the database, underlying cluster structure, and authentication methods can all change without significant changes needing to be made to `api.py`.

The most significant improvement that could be made to this structure would be to break the `UserManager` and `TranslationPodManager` instances out into separate programs running on their own Pods in the cluster. This would allow changes to be made in those classes, say to update the process of deleting Pods, without needing to update the entire Frontend/API Server Pod. For the time being, the relative simplicity of the API means that the added complexity and overhead of implementing a method to communicate between dedicated `TranslationPodManager` and `UserManager` Pods and Pods running the Frontend/API Server doesn't seem worthwhile. This

may change in the future.

Currently supported API requests are as follows:

1. `'/login'`: Logs in the user given by `'name'`, or, (given the rather lax development security) creates a user if none matching `'name'` exists. Returns the name of the user.
2. `'/channel'`: How users join RR Debug Sessions. Binds the user given by `'name'` to an RR Debug Session Pod matching `'channel'` in the database, if such a Pod exists. Returns the channel.
3. `'/pods'`: How the client gets a list of RR Debug Session Pods the user is currently participating in. Returns the channels for all RR Debug Session Pods who the user, given by `'name'`, is currently bound to in the database.
4. `'/examples'`: How the client gets a list of example RR Debug Session Pod names. Returns all names of example RR Debug Session Pods that exist in the database.
5. `'/new'`: Creates a new RR Debug Session Pod based on the name given by `'program'`. Binds the user, given by `'name'`, to the Pod. Returns the channel of the Pod. This request is the most complex, and is covered in more detail below.
6. `'/delete'`: Deletes a binding between the user given by `'name'` and the RR Debug Session Pod given by `'channel'`. If there are no bindings

left between the Pod and users (if all the users have left the debug session), deletes the Pod from the database and the cluster. Returns `True`.

All the API request handlers listed above return an error message in the case of an error occurring during the processing of an API request. Since most errors that could occur would either be the result of unanticipated issues on the part of the user, or of some unforeseen catastrophic internal error, informing users of error specifics would be unhelpful. Some `TranslationPodManager` and `UserManager` functions throw specific errors in the event of duplicate usernames, channels that do not exist, etc. These errors are caught and meaningful error messages are returned to the frontend webapp, which can then pass them on to the user.

To examine the process of processing an API request, it makes sense to look at the most complex example, creating a new RR Debug Session Pod. This example shows the process for processing an API POST request, interacting with the database, and dynamically creating a Pod. All API requests follow a similar procedure.

Processing a POST Request—Extracting Information Below is the route handler for the `/new` URL, as well as the instantiation of `TranslationPodManager` and `UserManager`.

```
tpm = podmanager.TranslationPodManager(  
    url='mongodb://database-load-balancer',
```

```

        port=27017)
um = userManager.UserManager(
        url='mongodb://database-load-balancer',
        port=27017)
app = Flask(__name__)
@app.route('/new', methods=['POST'])
def new():
    try:
        name = request.get_json()['name']
        program = request.get_json()['program']
        channel = tpm.create_pod([name], program)
        return {'channel': channel}
    except:
        return {'error': 'Internal Error'}

```

Listing 11: API Server New RR Debug Session Event Handler

All API routes only support POST requests. The route handler first extracts the relevant information from the POST request body, in this case `name` and `program`. `name` always corresponds to the user who is making the request's username, and `program` corresponds to the name of the example RR Debug Session image to be used.

`new` then calls `TranslationPodManager`'s `create_pod` function to create a new RR Debug Session Pod. `create_pod` can bind multiple users to a Pod when it is created, so `name` is passed inside a list.

The Database Before drilling into `create_pod`, it may be helpful to examine the MongoDB database which `TranslationPodManager` and `UserManager` interact with. The database consists of three collections: `users`, `pods`, and `examples`. Below is an example record from each collection:

`pods` and `users` have a many-to-many relationship. Each user can be bound to an unlimited number of pods, and vice versa. The `DatabaseController`

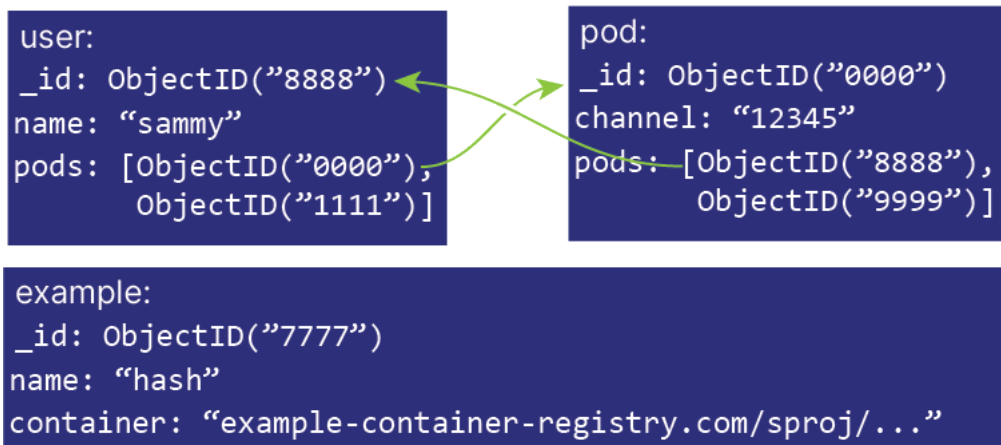


Figure 7: Database Record Examples

class includes functions to efficiently retrieve `users` and `pods` given one another, and to create/destroy bindings between `users` and `pods`.

`examples` are independent of `users` and `pods`. The container registry URLs of RR Debug Session example images are stored in the database so that new images can be easily added to the collaborative debugger upon creation.

MongoDB keeps `_ids` unique. In addition, the collaborative debugger defines uniqueness constraints on `users:name`, `pods:channel`, and `examples:name` when the database is created.

Processing a POST Request—Dynamic Pod Creation To dynamically create a Pod, `create_pod` gets the container registry URL of the image to base the Pod on, creates the new Pod, binds users to the Pod, and returns the channel the Pod was created with. The first part of the function is shown

below:

```
def create_pod(self, names, program):
    examples = self.get_examples()
    image = None
    try:
        image = examples[program]
    except:
        raise NoSuchExampleError
```

Listing 12: Pod Creation 1

`create_pod` calls `TranslationPodManager`'s `get_examples` method to get a list of RR Debug Session example image names and container registry URLs from the database. In the event that the user has somehow passed a spurious image name or has passed a name when there are no images, an exception is raised. Otherwise, a RR Debug Session Pod schema is created using the image:

```
dep = {'apiVersion': 'v1',
      'kind': 'Pod',
      'metadata': {'labels':
                    {'purpose': 'translate-rr'}},
      'spec': {'containers':
                [{'image': image,
                  'name': 'rr-test-container',
                  'command': ['sh'],
                  'args': ['startup.sh'],
                  'securityContext':
                    {'capabilities':
                     {'add': ['SYS_PTRACE']}}}]},
               'restartPolicy': 'OnFailure'}}
```

Listing 13: Pod Creation 2

This schema, represented as a dictionary in Python, is identical to the YAML representation of the RR Debug Session Pod creation schema shown in (4.1.2) with the exception of `args` which will be updated further in a later

step. Finally, `create_pod` does the work necessary to dynamically create a Pod:

```
channel = self.dbc.add_pod(
    self.dbc.get_userids_by_name(names))
dep['metadata']['name'] = channel
dep['spec']['containers'][0]['args'].append(channel)
resp = self.v1.create_namespaced_pod(
    body=dep, namespace='default')
# Wait to return the channel until the pod is live
# and read to receive incoming communication
status = self.v1.read_namespaced_pod_status(
    channel, 'default').status.container_statuses
while status == None or status[0].state.running == None:
    time.sleep(1)
    status = self.v1.read_namespaced_pod_status(
        channel, 'default').status.container_statuses
return channel
```

Listing 14: Pod Creation 3

First, `DatabaseController`'s `add_pod` method is called to insert a new `Pods` record into the database. `add_pod` randomly generates a new five digit `channel` for the Pod, using the uniqueness constraint imposed on `Pods:channel` to ensure an unused `channel` is generated. If no channels are open an error is raised. Otherwise, the users passed to `add_pod` are bound to the new Pod in the database, and `channel` is returned. This process has the effect of limiting the number of simultaneous RR Debug Sessions that can be in use at the same time to 89999, and of slowing as more channels are used. Given the few current users, the likelihood of a collision when generating a new channel is low enough that a more advanced method is unnecessary.

Next, the name of the container is updated to `channel`. In addition, the startup arguments are updated to include `channel`. This is how `channel` is

passed to the `startup.sh` script and eventually used to join a channel on the RR Message Server in `rrtranslation.py` (4.2.2 & 4.2.2).

The Kubernetes API is then finally used to create the Pod. The Pod is created in the default namespace. This could be changed to logically isolate dynamically and statically created Pods with minimal hassle, since Kubernetes' DNS resolution can connect pods accros namespaces. If Pod creation fails, an error will be thrown. Otherwise, a loop is used to wait until `rrtranslation.py` is running inside the Pod. Once the Pod is ready, `channel` is returned.

Processing a POST Request—Returning a Response Most request handlers, `'/new'` included, simply return a JSON serializable dictionary containing whatever results the API Server's instance of `TranslationPodManager` or `UserManager` returned, or an error message. The `'/new'` handler returns `'channel': channel` so that the frontend webapp's socket.io client can join the channel corresponding to the RR Debug Session Pod that was just created.

4.4 Webapp

TODO

5 Next Steps

TODO

References

- [1] Kubernetes Documentation.
- [2] Managed Kubernetes on DigitalOcean.
<https://www.digitalocean.com/products/kubernetes/>.
- [3] MongoDB. <https://www.mongodb.com>.
- [4] nginx. <https://nginx.org/en/>.
- [5] Yarn. <https://yarnpkg.com/en/>.
- [6] Flask, Nov. 2020. <https://github.com/pallets/flask>.
- [7] kubernetes/kubernetes, Nov. 2020. <https://github.com/kubernetes/kubernetes>.
- [8] Monaco editor, Nov. 2020. <https://github.com/microsoft/monaco-editor>.
- [9] Nodejs, Nov. 2020. <https://github.com/nodejs/node>.
- [10] React, Nov. 2020. <https://github.com/facebook/react>.
- [11] socket.io, Nov. 2020. <https://github.com/socketio/socket.io>.
- [12] xtermjs/xterm.js, Nov. 2020. <https://github.com/xtermjs/xterm.js>.
- [13] BROWN, N. C. C., SENTANCE, S., CRICK, T., AND HUMPHREYS, S.
Restart: The resurgence of computer science in uk schools. *ACM Trans. Comput. Educ.* 14, 2 (June 2014).

- [14] CHMIEL, R., AND LOUI, M. C. Debugging: From novice to expert. In *Proceedings of the 35th SIGCSE Technical Symposium on Computer Science Education* (New York, NY, USA, 2004), SIGCSE '04, Association for Computing Machinery, p. 17–21.
- [15] DOLAN-GAVITT, B., HODOSH, J., HULIN, P., LEEK, T., AND WHELAN, R. Repeatable reverse engineering with panda. In *Proceedings of the 5th Program Protection and Reverse Engineering Workshop* (New York, NY, USA, 2015), PPREW-5, Association for Computing Machinery.
- [16] FREE SOFTWARE FOUNDATION. *Debugging with GDB*, 2020.
- [17] GRINBERG, M. python-socketio, Nov. 2020. <https://github.com/miguelgrinberg/python-socketio>.
- [18] HANKS, B., MCDOWELL, C., DRAPER, D., AND KRNJAJIC, M. Program quality with pair programming in cs1. *SIGCSE Bull.* 36, 3 (June 2004), 176–180.
- [19] LI, C., CHAN, E., DENNY, P., LUXTON-REILLY, A., AND TEMPERO, E. Towards a framework for teaching debugging. In *Proceedings of the Twenty-First Australasian Computing Education Conference* (New York, NY, USA, 2019), ACE '19, Association for Computing Machinery, p. 79–86.

- [20] McCAULEY, R., FITZGERALD, S., LEWANDOWSKI, G., MURPHY, L., SIMON, B., THOMAS, L., AND ZANDER, C. Debugging: a review of the literature from an educational perspective. *Computer Science Education* 18, 2 (2008), 67–92.
- [21] McDOWELL, C., WERNER, L., BULLOCK, H. E., AND FERNALD, J. Pair programming improves student retention, confidence, and program quality. *Commun. ACM* 49, 8 (Aug. 2006), 90–95.
- [22] MICHAELI, T., AND ROMEIKE, R. Improving debugging skills in the classroom: The effects of teaching a systematic debugging process. In *Proceedings of the 14th Workshop in Primary and Secondary Computing Education* (New York, NY, USA, 2019), WiPSCE’19, Association for Computing Machinery.
- [23] O’CALLAHAN, R., JONES, C., FROYD, N., HUEY, K., NOLL, A., AND PARTUSH, N. Engineering record and replay for deployability: Extended technical report. *CoRR abs/1705.05937* (2017).
- [24] O’CALLAHAN, R., JONES, C., FROYD, N., HUEY, K., NOLL, A., AND PARTUSH, N. rr, Dec. 2019.
- [25] REPL.IT. The collaborative browser based IDE. <https://repl.it/>.
- [26] SMITH, C. pygdbmi, Nov. 2020. <https://github.com/cs01/pygdbmi>.