

The Development of a Collaborative Tool to Teach Debugging

A Senior Project submitted to
The Division of Science, Mathematics, and Computing
of
Bard College

by
Sammy Furr

Annandale-on-Hudson, New York
December, 2020

Abstract

Debugging is rarely formally taught, despite being used by programmers every day. Research indicates that it is valuable to teach debugging, and suggests that teaching it collaboratively may be maximally effective. The goal of this project is to create a collaborative debugger. The debugger aims to be the ideal platform to teach and learn debugging. This paper briefly reviews relevant literature covering teaching debugging and teaching programming collaboratively. Most of the paper is devoted to the design of the collaborative debugger.

1 Acknowledgments

- Keith O'Hara, my senior project advisor, for his constant thoughtfulness about my project, his advice about the future, and for giving me the opportunity to tutor for his classes.
- Mar Hicks and Kyle Hale, professors whose classes at IIT changed my perspective on computer science.
- All my professors at Bard, particularly Olga Voronina, Wakako Suzuki, and Nathan Shockey, whose encouragement of my interests outside of computer science has been instrumental in continuing my passion for it.
- My family and friends, for their constant love and support.

Contents

1	Acknowledgments	2
2	Introduction	6
2.1	Motivation	6
2.1.1	The Value of Teaching Debugging	7
2.1.2	Methods for Teaching Debugging	8
2.1.3	The Value of Collaborative Programming	9
3	Requirements	11
3.1	Encouraging Collaboration	11
3.2	Enhancing Student’s Understanding	12
3.3	Current Functionality	13
3.4	Considerations	20
4	Frontend Webapp	22
4.1	Tools Used	23
4.1.1	Mozilla’s rr	23
4.1.2	React	25
4.1.3	Monaco and Xterm.js	26
4.2	Overview	27
4.3	The Collaborative Interface to rr	29
4.3.1	Sending a Command	29

4.3.2	Receiving a Response	31
4.4	Considerations	33
5	Backend Design Overview	34
5.1	Tools Used	34
5.1.1	Docker	34
5.1.2	Kubernetes	34
5.2	Overview	35
5.3	Configuration and Setup	38
5.3.1	Cluster Requirements and Configuration	39
5.3.2	Creating Pods	39
5.3.3	Service Creation	43
6	RR Debug Sessions & RR Message Server	45
6.1	RR Message Server	47
6.1.1	Tools Used	47
6.1.2	Overview	47
6.2	RR Debug Session	49
6.2.1	Tools Used	49
6.2.2	Building Example Pods	50
6.2.3	The RR Translation Program	51
6.3	Considerations	54
7	Frontend/API Server	56

7.1	Tools Used	56
7.1.1	MongoDB	56
7.1.2	Flask and Node.js	56
7.2	Overview	57
7.3	Frontend/API Server Configuration	58
7.4	The API Server	59
7.4.1	Processing a POST Request—Extracting Information .	62
7.4.2	The Database	63
7.4.3	Processing a POST Request—Dynamic Pod Creation .	64
7.4.4	Processing a POST Request—Returning a Response . .	67
7.5	Considerations	67
8	Next Steps	68
8.1	Features to Add	68
8.2	Testing Efficacy	69
8.3	Conclusions	69

2 Introduction

Programmers spend large amounts of time debugging—the process of searching for and correcting errors in code that are not immediately apparent. They often use debuggers—tools which help programmers to inspect a running program—to assist them in this process of diagnosing sometimes inexplicable problems with their code. Though debuggers are primarily used to hunt down hidden mistakes, they are also powerful tools to understand program execution.

The goal of this project is to create a collaborative debugger to aid in the teaching of debugging. The collaborative debugger hopes to help teachers integrate this undertaught skill into their classes. It realizes the benefits of collaborative programming and teaching debugging by providing a platform that aims to be useful to both students and teachers.

2.1 Motivation

Debugging is invaluable in writing and understanding code, yet it is rarely formally taught [25]. Students are typically taught programming structures, concepts, and languages, but are left to learn the tools they use to write code alone. This approach often works well—a programmer’s choice of tools is often *very* personal and students figure out how to configure an individualized workflow. Perhaps because debuggers are tools, students are often expected to learn them with minimal guidance. Unlike editors or reference guides

however, effectively using a debugger requires a set of high-level, platform agnostic, teachable skills. Teaching these skills is effective, and translates into better and faster debugging and programming [24] [27]. Teaching these debugging skills collaboratively will likely offer the same confidence and correctness benefits realized by teaching programming collaboratively [23] [26].

2.1.1 The Value of Teaching Debugging

There is an unfortunate lack of research specifically into the efficacy of teaching debugging for computer science students, despite a recent rise in the inclusion of debugging in “computational thinking” curricula [27]. These curricula attempt to teach skills in computer science classes that translate into other subject areas: the UK’s computer science curriculum considers debugging an essential “transferable skill” [19].

There seems to be confidence that the problem-solving techniques used in debugging are widely applicable, but of greater interest to computer science teachers is whether teaching debugging directly benefits student programmers. Michaeli and Romeike conducted a good, albeit somewhat small, study on the efficacy of teaching a systematic debugging process to K12 students. They found that students who have been taught a specific debugging framework performed better in debugging tests and were more confident in their own debugging skills [27]. Their result is positive evidence towards the efficacy of teaching debugging, though it doesn’t include college or university students.

As Michaeli and Romeike point out, there is a lack of research into the value of teaching debugging in higher education. None of the research these authors found placed much focus on explicitly teaching debugging. Chmiel and Loui studied whether students who were provided with debugging tools and frameworks performed better on tests or spent less time on assignments than those who were not [20]. Though this research wasn't able to find conclusive evidence towards better performance on tests or assignments, it did find that students in the treatment group felt more confident in their debugging abilities. Unfortunately Chmiel and Loui's study didn't involve extended explicit teaching of debugging—use of the tools was voluntary, and variations in the students' individual abilities made the data difficult to evaluate.

Though there is a lack of higher-education research, the value of teaching debugging is still demonstrable. The research discussed all finds that K-12 and college students alike commonly resort to sporadic debugging techniques when beginning to learn. Since this pattern of behavior that explicitly teaching debugging corrects exists in college as well as in K-12 students, it seems logical that the benefit of explicitly teaching debugging to K-12 students should be realized equally by their collegiate counterparts.

2.1.2 Methods for Teaching Debugging

Similarly to research on the value of teaching debugging, research into how to best teach debugging is sparse. Chan et al. allow that “in general research on

how to improve debugging is sporadic”—an observation that leads them to research a framework to reduce the complexity of teaching debugging [24]. To organize their framework, they split debugging knowledge into 5 categories: *Domain*, *System*, *Procedural*, *Strategic*, and *Experiential*. They then review different debugging tools and teaching aids—from those that involve writing code to games—and map tools to the knowledge areas they seek to address. After an evaluation of a host of different tools, they claim to find a few significant faults in current debugging teaching platforms. The collaborative debugger primarily seeks to address the lack of back-tracing ability/coverage found in their research.

2.1.3 The Value of Collaborative Programming

Collaborative programming, where multiple programmers work together to write and test code, is popular in both industry and computer science education [26]. Collaborative programming most often manifests as pair programming—two programmers working together on the same program. Research into the value of pair programming is overwhelmingly positive. McDowell et. al. found that not only does pair programming significantly boost student confidence and the retention of students in computer science majors, but that it demonstrably improves student’s work [26]. These benefits of confidence and correctness are present when paired and non-paired students are given identical assignments [23], further indicating that the simple act of collaboration definitively benefits computer science students.

It seems that the benefits in confidence that result from teaching debugging should be magnified by teaching debugging collaboratively. There are similarities between the introductory nature of teaching fundamental debugging skills and the nature of teaching fundamental programming skills covered in the classes studied in [23] and [26]. Introductory programming classes typically use a specific programming language in order to introduce widely applicable programming principles. By using a specific collaborative platform to introduce debugging principles, students may realize the same benefits in both confidence *and* correctness that they do from pair programming in introductory computer science classes.

Debuggers exist at an intersection of tools and skills similar to programming languages themselves. By becoming familiar with a specific debugger, students may learn techniques and paradigms necessary to use all debuggers effectively. The collaborative debugger aims to provide the optimal platform for students to learn debugging skills.

3 Requirements

In order to provide the best platform for both teachers and students, the collaborative debugger must fulfill two key requirements:

1. It must encourage collaboration through a seamless experience.
2. It must make it easier to enhance students' understanding of debugging and programs being debugged.

3.1 Encouraging Collaboration

A large number of tools exist to facilitate collaborative programming. COVID-19 has greatly increased the demand for tools that not only make collaboration easier, but make it easy remotely. The tool that was most influential in the design of the collaborative debugger, Replit [15], enables remote collaborative programming.

Replit provides a simple browser-based IDE for over 50 languages. Programmers choose a language and can edit and run code collaboratively inside the Replit webapp. Similarly to other collaborative text-editors, input and output from all users is synced.

The collaborative debugger aims to provide a similar experience to tools like Replit. Students should be able to start a debugging session, easily invite other students, and interact with the debugger in a way that makes it easy to share their knowledge. By allowing multiple users to interact with the same

debugging session together in real-time, the collaborative debugger lets students share experience and work through problems together. The debugger also makes it easy for teachers to demonstrate debugging techniques remotely to multiple students, any of whom can also interact with the debugger. Since debugging sessions are hosted remotely, users can start a session in one location and resume it later from a different computer. This is particularly beneficial for students who may not always have access to a computer or a machine capable of running the programs they wish to debug.

3.2 Enhancing Student's Understanding

Apart from the benefits to understanding that are realized through collaboration, the collaborative debugger should enhance students' understanding of debugging in ways that a traditional debugger cannot. It aims to do this by allowing teachers to design and distribute lessons in the form of debugging sessions. Each debugging session is a lesson consisting of a deterministic recording of program execution, which students can debug repeatedly in order to build debugging skills and to learn about the execution process.

Currently, there are three example debug sessions included with the collaborative debugger:

1. **hash**: an example of classic pointer confusion. The hashtable implementation included with the program is slightly bugged, where failing to call `strncpy` results in the value for each key-value pair in the hash

pointing to the same location in memory. This is a particularly useful example for learning about setting watchpoints and breakpoints, as well as basic procedures to step through code.

2. **smash**: a brief example of stack smashing. A badly designed call to `sscanf` results in a portion of the call stack being overwritten, and the process is killed as a result of stack smashing. This example focuses on printing and examining memory locations.
3. **memoization**: a program that demonstrates the difference between a recursive, iterative, and memoized-recursive implementation of the classic `nth` Fibonacci number function. This example focuses on examining the stack under conditions produced by different types of functions.

The collaborative debugger should make it easy to distribute these sessions and for students to learn from them collaboratively.

3.3 Current Functionality

The following section details the first few steps students take to explore the **memoization** example using the current version of the collaborative debugger. This example, which focuses more on learning about how functions execute than on fixing bugs, highlights the ways in which the collaborative debugger enables pair programming.

Students A and B begin by examining the call stack at its deepest point in the recursive `nth` Fibonacci number function. The terminal interface to `rr`

and the source code view have been color-inverted to improve readability in this paper.

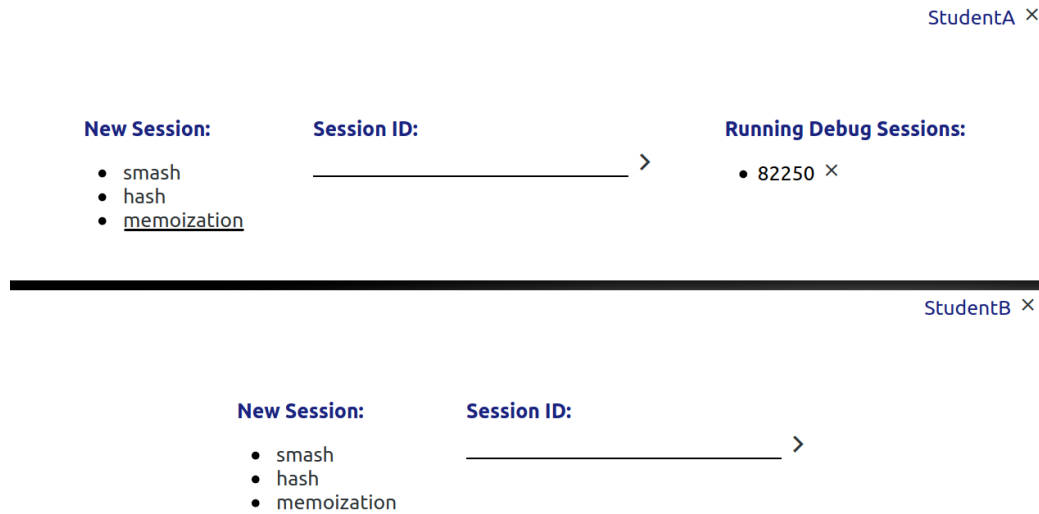


Figure 1: Creating a new Memoization Example Session

Students A and B both log into the collaborative debugger. To start a new session, Student A selects “memoization” from the list of example sessions.

StudentA ×

Session ID: 71825

memo.c

```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <time.h>
4
5 #define HASHSIZE 100
6
7 struct pair {
8     struct pair *next;
9     int key;
10    int value;
11 };
12
13 static struct pair *hashtable[HASHSIZE];
14
15 unsigned hash (int k)
16 {
17     return k % HASHSIZE;
18 }
19
20 struct pair *search(int k)
21 {
22     struct pair *pp;
23
24     for (pp = hashtable[hash(k)]; pp != NULL;
25          if (k == pp->key)
26              return pp;
```

rr

```
Welcome to Collab RR

(rr) run
0x00007f407851a103 in ?? () from /lib64/ld-linux-x86-64.so.2
The program being debugged has been started already.
Start it from the beginning? (y or n) [answered Y; input not from terminal]
Starting program: /home/debug/.local/share/rr/memo-0/mmap_hardlink_3_memo
Program stopped.
0x00007f407851a100 in ?? () from /lib64/ld-linux-x86-64.so.2

(rr) □
```

StudentB ×

New Session:

- smash
- hash
- memoization

Session ID:

71825 >

Figure 2: Joining the Memoization Example Session

In order to join Student A’s session, Student B uses session code “71825”. Meanwhile, Student A is presented with the the main view of the collaborative debugger. On the left is a view of the source code for the **memoization** example. On the right is the terminal interface to **rr**, the enhancement to **gdb** used by the collaborative debugger. Students debug recordings of program execution, which can be replayed deterministically. **rr** is covered in detail in 4.1.1.

StudentA ✕

Session ID: 71825

memo.c

```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <time.h>
4
5 #define HASHSIZE 100
6
7 struct pair {
8     struct pair *next;
9     int key;
10    int value;
11 };
12
13 static struct pair *hashtable[HASHSIZE];
14
15 unsigned hash (int k)
16 {
17     return k % HASHSIZE;
18 }
19
20 struct pair *search(int k)
21 {
22     struct pair *pp;
23
24     for (pp = hashtable[hash(k)]; pp != NULL;
25         if (k == pp->key)
26             return pp;
```

rr

```
Welcome to Collab RR

(rr) run
0x00007f407851a103 in ?? () from /lib64/ld-linux-x86-64.so.2
The program being debugged has been started already.
Start it from the beginning? (y or n) [answered Y; input not from terminal]
Starting program: /home/debug/.local/share/rr/memo-0/mmap_hardlink_3_memo
Program stopped.
0x00007f407851a100 in ?? () from /lib64/ld-linux-x86-64.so.2

(rr) c
Continuing.
Fib recursive: time=0.009810 result=832040
Fib iterative: time=0.000002 result=832040
Fib memoization: time=0.000004 result=832040
Program received signal SIGKILL, Killed.
0x0000000070000002 in ?? ()

(rr) □
```

StudentB ✕

Session ID: 71825

memo.c

```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <time.h>
4
5 #define HASHSIZE 100
6
7 struct pair {
8     struct pair *next;
9     int key;
10    int value;
11 };
12
13 static struct pair *hashtable[HASHSIZE];
14
15 unsigned hash (int k)
16 {
17     return k % HASHSIZE;
18 }
19
20 struct pair *search(int k)
21 {
22     struct pair *pp;
23
24     for (pp = hashtable[hash(k)]; pp != NULL;
25         if (k == pp->key)
26             return pp;
```

rr

```
Welcome to Collab RR

(rr) c
Continuing.
Fib recursive: time=0.009810 result=832040
Fib iterative: time=0.000002 result=832040
Fib memoization: time=0.000004 result=832040
Program received signal SIGKILL, Killed.
0x0000000070000002 in ?? ()

(rr) █
```

Figure 3: Observing Output from the Memoization Example Session

Student B now types `c` to continue execution. Both students see the command and output: execution times and results for three different implementations of `nth Fibonacci`. An advantage of `rr`'s recording model is that

these numbers will be the same every time the students run the program.

StudentA ✕

Session ID: 71825

memo.c

```

40 hashtable = hashtable;
42 pp->next = hashtable[hashval];
43 hashtable[hashval] = pp;
44 }
45 pp->value = value;
46 return pp;
47 }
48 int fib_rec (int n){
49     if (n < 0)
50         return -1;
51     if (n == 0 || n == 1)
52         return n;
53     return fib_rec(n - 1) + fib_rec (n - 2);
54 }
55
56 int fib_itr (int n) {
57     int a = 0;
58     int b = 1;
59     for (int x = 0; x < n; x++){
60         int t = a;
61         a = b;
62         b = t + b;
63     }
64     return a;
65 }

```

rr

```

Fib iterative: time=0.000002 result=832040
Fib memoization: time=0.000004 result=832040
Program received signal SIGKILL, Killed.
0x0000000070000002 in ?? ()

(rr) break 52
Breakpoint 1 at 0x557268dbf303: file memo.c, line 52.

(rr) run
[Inferior 1 (process 16) exited normally]
Starting program: /home/debug/.local/share/rr/memo-0/mmap_hardlink_3_memo
Program stopped.
0x00007f407851a100 in ?? () from /lib64/ld-linux-x86-64.so.2

(rr) c
Continuing.
Breakpoint 1, fib_rec (n=1) at memo.c:52
52     return n;

(rr)

```

StudentB ✕

Session ID: 71825

memo.c

```

1  #include <stdlib.h>
2  #include <stdio.h>
3  #include <time.h>
4
5  #define HASHSIZE 100
6
7  struct pair {
8      struct pair *next;
9      int key;
10     int value;
11 };
12
13 static struct pair *hashtable[HASHSIZE];
14
15 unsigned hash (int k)
16 {
17     return k % HASHSIZE;
18 }
19
20 struct pair *search(int k)
21 {
22     struct pair *pp;
23
24     for (pp = hashtable[hash(k)]; pp != NULL;
25         if (k == pp->key)
26             return pp;

```

rr

```

Fib iterative: time=0.000002 result=832040
Fib memoization: time=0.000004 result=832040
Program received signal SIGKILL, Killed.
0x0000000070000002 in ?? ()

(rr) break 52
Breakpoint 1 at 0x557268dbf303: file memo.c, line 52.

(rr) run
[Inferior 1 (process 16) exited normally]
Starting program: /home/debug/.local/share/rr/memo-0/mmap_hardlink_3_memo
Program stopped.
0x00007f407851a100 in ?? () from /lib64/ld-linux-x86-64.so.2

(rr) c
Continuing.
Breakpoint 1, fib_rec (n=1) at memo.c:52
52     return n;

(rr)

```

Figure 4: Setting a Breakpoint in the Memoization Example Session

Student A enters `break 52` to set a breakpoint at line 52, the base case of `fib_rec`. Student B types `run` and then `c` to restart the recording and continue to the breakpoint.

Session ID: 71825

```

memo.c
40 hashtable = hashtable;
41 pp->next = hashtable[hashval];
42 hashtable[hashval] = pp;
43 }
44 pp->value = value;
45 return pp;
46 }
47
48 int fib_rec (int n){
49     if (n < 0)
50         return -1;
51     if (n == 0 || n == 1)
52         return n;
53     return fib_rec(n - 1) + fib_rec (n - 2);
54 }
55
56 int fib_itr (int n) {
57     int a = 0;
58     int b = 1;
59     for (int x = 0; x < n; x++){
60         int t = a;
61         a = b;
62         b = t + b;
63     }
64     return a;
65 }
66

rr
#12 0x0000557268dbf315 in fib_rec (n=13) at memo.c:53
#13 0x0000557268dbf315 in fib_rec (n=14) at memo.c:53
#14 0x0000557268dbf315 in fib_rec (n=15) at memo.c:53
#15 0x0000557268dbf315 in fib_rec (n=16) at memo.c:53
#16 0x0000557268dbf315 in fib_rec (n=17) at memo.c:53
#17 0x0000557268dbf315 in fib_rec (n=18) at memo.c:53
#18 0x0000557268dbf315 in fib_rec (n=19) at memo.c:53
#19 0x0000557268dbf315 in fib_rec (n=20) at memo.c:53
#20 0x0000557268dbf315 in fib_rec (n=21) at memo.c:53
#21 0x0000557268dbf315 in fib_rec (n=22) at memo.c:53
#22 0x0000557268dbf315 in fib_rec (n=23) at memo.c:53
#23 0x0000557268dbf315 in fib_rec (n=24) at memo.c:53
#24 0x0000557268dbf315 in fib_rec (n=25) at memo.c:53
#25 0x0000557268dbf315 in fib_rec (n=26) at memo.c:53
#26 0x0000557268dbf315 in fib_rec (n=27) at memo.c:53
#27 0x0000557268dbf315 in fib_rec (n=28) at memo.c:53
#28 0x0000557268dbf315 in fib_rec (n=29) at memo.c:53
#29 0x0000557268dbf315 in fib_rec (n=30) at memo.c:53
#30 0x0000557268dbf426 in timer (f=0x557268dbf2da <fib_rec>, i=30, result=0x7ffe99602d80) at memo.c:82
#31 0x0000557268dbf487 in main () at memo.c:91

```

Session ID: 71825

```

memo.c
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <time.h>
4
5 #define HASHSIZE 100
6
7 struct pair {
8     struct pair *next;
9     int key;
10    int value;
11 };
12
13 static struct pair *hashtable[HASHSIZE];
14
15 unsigned hash (int k)
16 {
17     return k % HASHSIZE;
18 }
19
20 struct pair *search(int k)
21 {
22     struct pair *pp;
23
24     for (pp = hashtable[hash(k)]; pp != NULL;
25          if (k == pp->key)
26              return pp;

```

```

rr
(rr) bt
#0 fib_rec (n=1) at memo.c:52
#1 0x0000557268dbf315 in fib_rec (n=2) at memo.c:53
#2 0x0000557268dbf315 in fib_rec (n=3) at memo.c:53
#3 0x0000557268dbf315 in fib_rec (n=4) at memo.c:53
#4 0x0000557268dbf315 in fib_rec (n=5) at memo.c:53
#5 0x0000557268dbf315 in fib_rec (n=6) at memo.c:53
#6 0x0000557268dbf315 in fib_rec (n=7) at memo.c:53
#7 0x0000557268dbf315 in fib_rec (n=8) at memo.c:53
#8 0x0000557268dbf315 in fib_rec (n=9) at memo.c:53
#9 0x0000557268dbf315 in fib_rec (n=10) at memo.c:53
#10 0x0000557268dbf315 in fib_rec (n=11) at memo.c:53
#11 0x0000557268dbf315 in fib_rec (n=12) at memo.c:53
#12 0x0000557268dbf315 in fib_rec (n=13) at memo.c:53
#13 0x0000557268dbf315 in fib_rec (n=14) at memo.c:53
#14 0x0000557268dbf315 in fib_rec (n=15) at memo.c:53
#15 0x0000557268dbf315 in fib_rec (n=16) at memo.c:53
#16 0x0000557268dbf315 in fib_rec (n=17) at memo.c:53
#17 0x0000557268dbf315 in fib_rec (n=18) at memo.c:53
#18 0x0000557268dbf315 in fib_rec (n=19) at memo.c:53
#19 0x0000557268dbf315 in fib_rec (n=20) at memo.c:53
#20 0x0000557268dbf315 in fib_rec (n=21) at memo.c:53

```

Figure 5: Examining the Stack in the Memoization Example Session

The students examine the stack resulting from a call of `fib_rec(30)`. Students can inspect different areas of the program and output while collaborating—Student A can look at `fib_rec`, while Student B scrolls up for a preview of the memoized implementation.

The students will continue to inspect the stack for other functions, exploring the reasons for the vast performance disparity between `fib_memo` and `fib_rec`. Students A and B can not only help each other to learn about the stack and memoization, but can also teach one another how to use a debugger. Because the frontend produces near-identical output to `rr`, students can directly translate their knowledge to it or `gdb` in the future. Paradigms, such as examining the stack to understand execution, can be expanded to all future debugging.

3.4 Considerations

Though the collaborative debugger makes it easy for students in the above example to practice using a debugger together, more features could be added to help improve students' understanding of the program being debugged. Students using all three examples would benefit greatly from visualization tools to help understand pointer locations, the process space, and data structures.

In `memoization`, a simple stack visualization that updated with each command sent to the debugger would allow students to more easily see changes in the process space. Rather than having to repeatedly `continue` and `backtrace` to visualize the stack, students could simply step through relevant portions of the program and watch as the stack updated. This sort of visualization pairs particularly well with `reverse-continue` or `reverse-step(i)`, allowing students to move back and forth over particularly significant stack changes. Visualization tools also aid students' understanding of debugging—

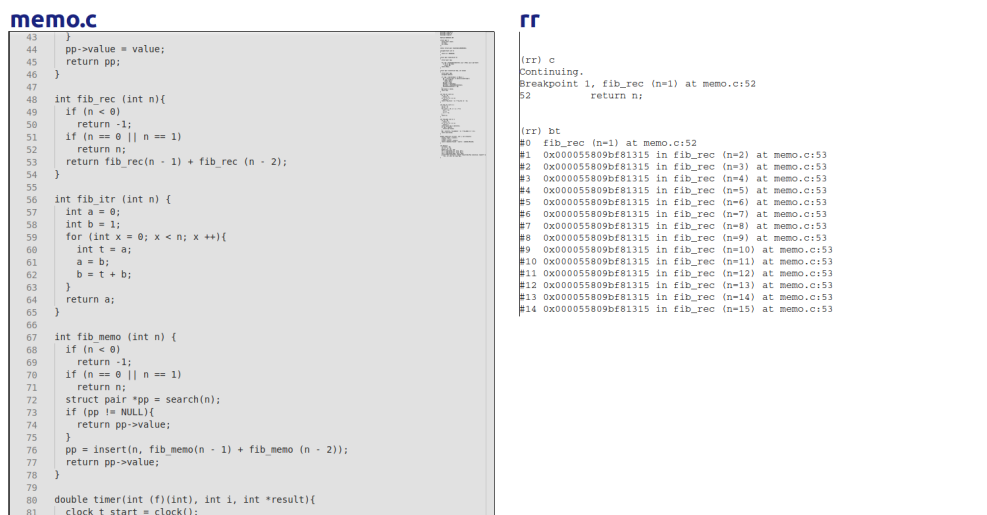
students can see how typing `bt` to backtrace or `info frame` to view information about the current stack frame provide the same information as the visualization tools they are already familiar with. As discussed in 4.4, adding visualization tools to the collaborative debugger should be relatively easy.

The aim of the collaborative debugger is to provide students with an experience they can directly reference when using similar debuggers in the future, as well as one which is seamless to encourage the acquisition of general debugging skills. Creating such a platform will hopefully encourage both students and teachers to explore the undertaught area of debugging.

4 Frontend Webapp

An overview of the `memoization` example discussed in 3.3 is below. This pattern is followed by all debugging sessions in the frontend of the collaborative debugger.

Session ID: 17734



The screenshot displays a collaborative debugger interface. On the left, a code editor shows the source code for `memo.c`. The code includes a recursive `fib_rec` function, an iterative `fib_itr` function, and a memoized version `fib_memo` that uses a hash table for caching. On the right, a memory dump window titled `rr` shows the state of memory at a specific point, listing addresses and values for various variables and function calls.

```
memo.c
43 }
44 pp->value = value;
45 return pp;
46 }
47
48 int fib_rec (int n){
49     if (n < 0)
50         return -1;
51     if (n == 0 || n == 1)
52         return n;
53     return fib_rec(n - 1) + fib_rec (n - 2);
54 }
55
56 int fib_itr (int n) {
57     int a = 0;
58     int b = 1;
59     for (int x = 0; x < n; x++){
60         int t = a;
61         a = b;
62         b = t + b;
63     }
64     return a;
65 }
66
67 int fib_memo (int n) {
68     if (n < 0)
69         return -1;
70     if (n == 0 || n == 1)
71         return n;
72     struct pair *pp = search(n);
73     if (pp != NULL){
74         return pp->value;
75     }
76     pp = insert(n, fib_memo(n - 1) + fib_memo (n - 2));
77     return pp->value;
78 }
79
80 double timer(int (f)(int), int i, int *result){
81     clock_t start = clock();
```

```
rr
(rr) c
Continuing.
Breakpoint 1, fib_rec (n=1) at memo.c:52
52     return n;

(rr) bt
#0 fib_rec (n=1) at memo.c:52
#1 0x000055809bf81315 in fib_rec (n=2) at memo.c:53
#2 0x000055809bf81315 in fib_rec (n=3) at memo.c:53
#3 0x000055809bf81315 in fib_rec (n=4) at memo.c:53
#4 0x000055809bf81315 in fib_rec (n=5) at memo.c:53
#5 0x000055809bf81315 in fib_rec (n=6) at memo.c:53
#6 0x000055809bf81315 in fib_rec (n=7) at memo.c:53
#7 0x000055809bf81315 in fib_rec (n=8) at memo.c:53
#8 0x000055809bf81315 in fib_rec (n=9) at memo.c:53
#9 0x000055809bf81315 in fib_rec (n=10) at memo.c:53
#10 0x000055809bf81315 in fib_rec (n=11) at memo.c:53
#11 0x000055809bf81315 in fib_rec (n=12) at memo.c:53
#12 0x000055809bf81315 in fib_rec (n=13) at memo.c:53
#13 0x000055809bf81315 in fib_rec (n=14) at memo.c:53
#14 0x000055809bf81315 in fib_rec (n=15) at memo.c:53
```

Figure 6: Debugging the Memoization Example Session

1. Students A and B both log into the collaborative debugger from different computers.
2. Student A creates a new Debug Session by clicking the button for `memoization` under “New Session”. A few moments later, they’re presented with the main screen of the debugger.

3. Using the session code (17734 in the example above), Student B joins the Debug Session.
4. Both students can now interact with the same debugger remotely. Input and output from the debugger and the current source file are synced between all students. Students can work together to understand the reasons why memoization often improves the execution time of recursive programs.

4.1 Tools Used

The frontend of the collaborative debugger must handle constant changes in state as users participate in the process outlined above. The frontend uses React, Monaco, and Xterm.js to manage state and display output. It primarily displays the output from `rr`, an enhancement of `gdb` that the collaborative debugger uses to enable collaboration.

4.1.1 Mozilla’s `rr`

Notice that the debugging window in the above example is called “`rr`”. `rr` is “a lightweight tool for recording, replaying and debugging execution of applications” [29]. `rr` allows a programmer to record the execution of a program on any compatible machine and replay the execution later. This enhances GDB’s ability to “time-travel” when debugging, using commands such as `reverse-continue` and `reverse-step(i)` [22] to step backwards and for-

wards through a program’s execution. Through a novel encapsulation of the execution space, rr is able to deterministically record and replay the execution of syscalls and other process behavior that differs run-to-run. This is invaluable when trying to debug behavior that is not entirely dependent on the code being debugged. A typical workflow in rr consists of recording an inexplicable error, replaying execution to find the area in which the error occurs, and then narrowing in on the bug not by re-running the entire program, but by progressing back and forth through execution in the problem area.

rr is an ideal tool for teaching debugging because it allows instructors to record execution of a program and design a debugging example with the knowledge that normally non-deterministic events will be repeatable, and that any input they provide to the program will be exactly replicated. With the collaborative debugger, teachers can record a program’s execution and design a debugging lesson which students can work on together. The repeatability of rr means that students can focus on debugging, and teachers can create as specific examples as they please. The use of rr is the most significant step the collaborative debugger takes to addressing the lack of back-tracing ability/coverage found in existing tools for teaching debugging [24].

In comparison to solutions like PANDA [21] that rely on capturing the entire state of a virtual machine to replay execution, rr records and replays faster, produces far smaller files, and doesn’t force execution inside of a VM. [28] The trade off for these benefits are two major system limitations: rr is only compatible with the Linux kernel, and it’s deterministic recording

and replay relies on a feature that is only found on modern *Intel* x86 CPUs. These limitations influenced the development of this project as a webapp similar to existing tools for collaborative programming.

Luckily, the speed and size benefits of rr lend themselves well to non-local execution. In conjunction with tools used to create the backend of the collaborative debugger it takes a few seconds to create a new debug session running rr and connect it to the frontend webapp discussed in this section.

4.1.2 React

React is JavaScript library that simplifies creating user interfaces and managing state [12]. React’s state management is of particular importance to the collaborative debugger’s frontend. State constantly changes as users create/delete debugging sessions, join existing sessions, and communicate with rr. React allows classes and functions to encapsulate components such as a list of existing debugging sessions, a view of the current program’s source code, and the terminal interface with rr. Instances of these classes maintain state and update efficiently.

The frontend makes extensive use of JSX, syntax which allows the inclusion of segments of HTML code within a React app written in JavaScript. This makes it easy for each component of the one-page webapp to hide/show subcomponents as state changes.

4.1.3 Monaco and Xterm.js

After joining or creating a debugging session, users spend most of their time interacting collaboratively with rr. Their primary interface to rr is through Xterm.js, a frontend component that makes it easy to emulate terminal behavior in the browser [17]. With a few control methods, it is simple to provide a terminal interface to rr that is virtually indistinguishable from a local session. By using the Xterm.js based interface, students can learn to use rr (and by extension gdb) collaboratively, and directly translate that knowledge to individual work.

In addition to the terminal interface, the frontend shows a view of the current source file being debugged. The Monaco Editor [6] is used to display this source view. Though more complex than is strictly necessary to display code, Monaco makes it easy to format and syntax-highlight. Using Monaco also simplifies the future addition of editing source code, should the need arise. React’s state management allows updating text in the editor as efficiently as possible.

In order to easily integrate Monaco and Xterm.js into the React frontend, React wrapper libraries were used. These are react-xterm [14] and react-monaco-editor [13].

4.2 Overview

React makes it easy to create simple webapps with no build configuration necessary through the use of the `create-react-app` command. This creates a structure and toolchain for a basic one-page application—suitable for the collaborative debugger. Once this structure was created and any superfluous files/code were removed, it wasn't necessary for the purposes of the collaborative debugger to directly edit any files except two: `App.js` and `App.css`. The pre-defined build process takes care of transforming the React code defined in `App.js` and the styles defined in `App.css` into portable JavaScript, HTML, and CSS.

`App.js` consists of a primary React Component class, `Debugger`, as well as many secondary classes and functions. More complex components of the frontend, such as the `Xterm.js` interface to `rr`, are created as classes, while simple components like buttons are defined as functions. All classes contain a `render()` method responsible for returning HTML elements to display based on the state contained in the class. `Debugger`'s `render()` method is responsible for displaying all other components, based on whether the user is logged in, has selected a new session, etc.

Many components “lift state up” to `Debugger`. This is accomplished by passing methods from `Debugger` in as “props” when elements are created in `Debugger`'s `render()` method:

```
onLogin(name) {
```

```
    this.setState({user: name})
    :
  }
  // In Debugger.render(). This is JSX.
  login = <LoginForm onLogin={this.onLogin}/>;
  // In LoginForm.
  this.props.onLogin(data.name);
```

Listing 1: Lifting State Up

Debugger’s `onLogin(name)` method changes the state of debugger to indicate that the user is logged in. When an instance of `LoginForm` is created in `Debugger.render()`, `onLogin(name)` is passed to it as a “prop”. The code for logging in the user can be contained in `LoginForm`, but when it becomes necessary to change the state of the whole application to indicate that the user is logged in, `LoginForm` calls `Debugger.onLogin(name)`, lifting state up to `Debugger`.

Other than methods to handle state and display components, `Debugger` also contains the `Socket.IO` client that is used to communicate with the actual debugger running in the collaborative debugger’s backend. `Socket.IO` is a library that extends `WebSockets`. It and this communication process will be discussed in detail in (6.1). Similarly to how `onLogin(name)` is passed as a property to `LoginForm`, the `Socket.IO` client is passed to the most important component in the frontend, `RRTerm`, when it is instantiated.

4.3 The Collaborative Interface to rr

As shown in Figure 6, the primary piece of the collaborative debugger’s frontend is an interface to the rr debugger backend. This interface currently consists of a view of the user’s source code and a terminal interface to rr. These are contained within the most important class in the frontend: `RRTerm`. This section will outline the process of sending an rr command to the backend and displaying a response. Apart from sending simple API requests to start sessions or log users in and out, this communication process is where the frontend webapp spends most of its time.

In debugging the `memoization` example program, users will want to display the call stack from within each implementation of the `nth` Fibonacci number function. To stop in `fib_rec`, the recursive implementation, one user will want to set a breakpoint at line 52 in the source file. To do this they will type the command `break 52` and hit enter. The command will be sent to the backend, and all users in the same debug session (17734 in Figure 6) will see the debuggers response:

```
Breakpoint 1 at 0x561c8fede303: file memo.c, line 52.
```

Users who did not send the command `break 52` will see it displayed along with the response.

4.3.1 Sending a Command

`RRTerm`’s state is instantiated as such:

```
this.state = {command: '',
              name:  '',
              linum: 0,
              code: ''};
```

Listing 2: RRTerm's State

`name`, `linum`, and `code` store information about the current source file being debugged. `command` is the most important field for examining how commands are sent to the backend, as it holds the current command. *Every time* the user types a character, `command` is updated to reflect the current state of the Xterm.js component. The code to handle keystrokes is shown below:

```
this.term.on('key', (key, ev) => {
  const printable = !ev.altKey && !ev.ctrlKey && !ev.metaKey;
  if (ev.key === 'Enter') {
    // Disable input while we wait for the response from rr
    this.term.disableStdin = true;
    const c = this.state.command;
    this.sendCommand(c);
    this.setState({command: ''});
  } else if (ev.key === 'Backspace') {
    // Do not delete the prompt
    if (this.term.buffer._activeBuffer.x > prompt_length) {
      this.term.write('\b \b');
      this.updateCommand(this.state.command, 'delete');
    }
  } else if (ev.key === 'ArrowUp') {
  } else if (ev.key === 'ArrowDown') {
  } else if (printable) {
    this.term.write(key);
    this.updateCommand(this.state.command, key);
  }
});
```

Listing 3: Handling Keystrokes

Whenever the user types a printable key, the key is written to the terminal and `command` is updated using the `updateCommand` method. If a user types

`brk`, `updateCommand` would be called three times, once for each keystroke. Upon realizing they made a mistake when typing `break 52`, the user would likely delete a character. As long as the prompt (`(rr)` followed by a space) would not be overwritten, the terminal is updated and `updateCommand` is called again. When the user hits the enter key the command is finally sent to the backend using `sendCommand`. To mimic the behavior of a local `rr` debug session and to prevent garbled input, writing to the terminal is disabled until a response is received.

4.3.2 Receiving a Response

All users connected to the same debug session will receive responses to the commands sent by one user. This is how collaboration is achieved in the collaborative debugger. The JSON response object to `break 52` received by all users is shown below:

```
{
  channel: '11734',
  command: 'break 52',
  from: '1ed644ad87d44cd68121de3daa65adb4',
  response: {
    output: 'Breakpoint 2 at 0x561c8fede303: file memo.c, line 52.'
    source: {
      file_name: null,
      current_line: '87',
      contents: null
    }
  }
}
```

Listing 4: JSON Response

The command has been issued from a breakpoint set at `main`, which is at

line 87 in `memo.c`. Since the source file hasn't changed, `source.file_name` and `source.contents` are set to `null`. `output` contains the response from `rr`, while `channel`, `command`, and `from` contain information necessary to display commands correctly for all users.

```
this.props.socket.on('rr_response', (data) => {
  if(data.from !== this.props.socket.id){
    // Display commands from other clients
    this.term.disableStdin = true;
    this.term.clear();
    this.term.write(data.command);
  };
  // Write a newline, but no prompt
  this.term.writeln('');
  // Write each line of the output to the terminal
  data.response.output.split(/\n+/).forEach((l) => {
    this.term.writeln(l);
  });
  this.term.prompt();
  // Re-enable input. See above and below.
  this.term.disableStdin = false;
  // If an error has occurred, there won't be any new
  // information about the location in the source file.
  if(data.response.source !== null) {
    this.updateSource(data.response.source);
  }
});
```

Listing 5: Processing a Response

When a response is received, `RRTerm` first checks to see if the response has originated from a different user. If so, input is disabled, the current line is cleared, and the command that corresponds to the response is written to the terminal. Each line of the response is then written to the terminal. Finally, input is reenabled, having been disabled either to write commands received from another user, or when the command was first sent. `updateSource` is called to update the Monaco editor source file view. Since `break 52` was sent

from a location already inside `memo.c`, the source file name and contents are not updated, though the line number is changed.

4.4 Considerations

The intention of much of the code described above, such as that to disable input after a command has been sent, is to mimic `rr` in a way that is invisible to the user. The user experience is intended to be that of using a fully-fledged debugger collaboratively. The source code view allows users to easily see what code they are working on together, and the terminal updates quickly whenever a user enters a command.

Care has been taken to design the frontend in a modular fashion so that it can be easily extended in the future. Visualization tools could be easily added by sending additional data about program state in the response and creating new components. If further authentication features or the ability to edit source code are desired in the future, relevant components can be updated without the need to change the entire application. Since the collaborative debugger depends heavily on interactions between its frontend and backend, similar care has been taken in designing the backend, which will be described in the following sections.

5 Backend Design Overview

The design of the collaborative debugger’s backend is heavily distributed, allowing individual components to be modified without the whole system needing to be reconfigured. Kubernetes in conjunction with Docker are used to create the backend of the collaborative debugger.

5.1 Tools Used

Kubernetes and Docker were chosen largely so that new Debug Sessions could be easily and securely created on demand.

5.1.1 Docker

Docker is containerization software. Docker containers encapsulate applications and their dependencies while sharing their host’s kernel. This allows for lightweight, portable, and secure execution of complex programs [1]. All programs in the collaborative debugger’s backend run in Docker containers. Docker containers’ lightweight nature and security are instrumental when trying to quickly create containers to run user-defined code.

5.1.2 Kubernetes

Kubernetes is the de facto standard in container orchestration software. It provides a layer of abstraction on top of normal containers, like those created by Docker. By bundling one or more closely linked containers into a “Pod”,

Kubernetes is able to manage deployment and re-deployment of applications running inside containers. It is trivial to create new Pods (or multiple Pods running the same application) as needed within a Kubernetes cluster [3]. The speed at which even relatively large Pods can be created and the inherent security provided by containerization drove the decision to create a new Pod on the fly for each debugging instance in the collaborative debugger. This allows the debugger to provide a similar level of convenience to existing collaborative tools, such as Replit (3.1).

Kubernetes also provides services to facilitate load balancing, manage storage volumes, and contain secrets. The abstraction provided by these features, in tandem with the ease of Kubernetes deployment on a managed Kubernetes service [5] greatly accelerated development.

5.2 Overview

The collaborative debugger consists of a distributed Kubernetes backend and frontend React webapp. Kubernetes was chosen for the backend primarily so that a Pod could be created dynamically for each debugging session.

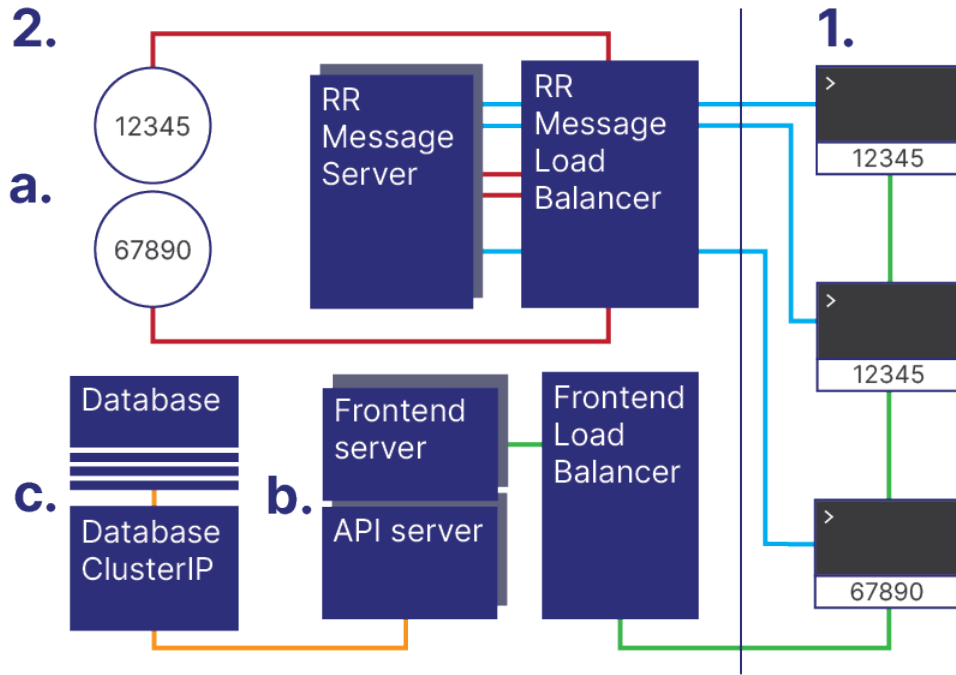


Figure 7: Detailed Overview of the Collaborative Debugger

Each backend component of the collaborative debugger runs in it's own individual Pod. There are two different classes of Pods in the collaborative debugger:

1. Statically created Pods. These are the *RR Message Server Pods*, the *Frontend/API Server Pods*, and the *Database Pod*. This class of Pods are manually created when the cluster that will run the backend is first initialized. The RR Message Server Pods and Frontend/API Server Pods may be created using Deployments [4] to allow scaling, where multiple Pods running the same application may be created to facilitate increased load.

2. Dynamically created Pods. This class of Pod contains the individual instances of *RR Debug Session Pods* that are created on request by the API server. When a client requests a new debug session, the API server uses the Kubernetes API to create a new Pod based on an existing template, gives the Pod a unique identifier, and associates it in the database with the requesting client.

These Pods communicate with other Pods in the cluster and with the outside world through Services. The Kubernetes documentation defines a Service as “an abstraction which defines a logical set of Pods and a policy by which to access them” [4]. In the collaborative debugger, these Services manifest as:

1. The *Database ClusterIP*: a ClusterIP, which exposes the Database Pod only inside the cluster. The only component that makes use of this ClusterIP is the API server, which uses it primarily to communicate information about users and RR Debug Sessions with the database.
2. The *RR Message Load Balancer*: a Load Balancer, which exposes the RR Message Server Pods to the outside world. Using Socket.IO, clients send commands to and receive responses from individual RR Debug Session Pods through the RR Message Server Pods.
3. The *Frontend Load Balancer*: another Load Balancer, which exposes the Frontend/API Server Pods to the outside world. Clients request the

frontend webapp and send API requests/receive API responses through this Load Balancer.

The frontend webapp dynamically updates as the user requests new debug sessions, issues commands to rr, and visits new source files. A user can be part of multiple debug sessions simultaneously. Each debug session is assigned at 5 digit identifier at creation, which is used to join sessions in progress.

Each component of the backend will be discussed in depth in the following sections.

5.3 Configuration and Setup

Configuration and setup of the collaborative debugger is relatively simple. After a Kubernetes cluster is created (this is made easier by using a Managed Kubernetes service) Pods and Services are created using various configuration files. Services should be created first, so that Pods which rely on access to Services function properly on creation. The following sections outline the process of creating a cluster, Pods, and Services, with a focus on statically created Pods. The process of building images for building dynamically created Pods is similar, but the process of starting the Pods is more complicated. This process will be discussed in-depth in the section on the API server (7).

5.3.1 Cluster Requirements and Configuration

Though Kubernetes aims to be largely platform-agnostic, the requirements of rr impose some restrictions on cluster setup and configuration. Clusters, even those running inside a VM, must be run on machines using relatively modern Intel x86 CPUs (Nehalem and beyond). The clusters must run on an operating system using Linux kernel version 3.11 or higher [29]. Finally, in order for rr to be able to work efficiently, the `kernel.perf_event_paranoid` parameter must be set to 1 [29]. This should be done on every node in the cluster which will run RR Debug Sessions. For the purposes of development, it has been set to 1 on all nodes in the collaborative debugger cluster.

5.3.2 Creating Pods

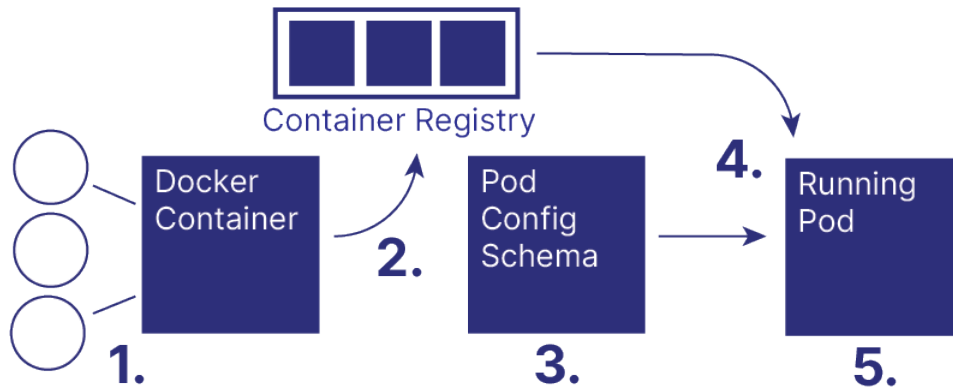


Figure 8: The Pod Creation Process

Pods are created in five steps:

1. A `Dockerfile` is used to build a new Docker container image from

various pieces of source code, scripts, and a base image (such as the official MongoDB image or official Ubuntu image). The Dockerfile also contains instructions to install necessary packages, run build scripts, and create file structures in the image.

2. The Docker image is tagged and uploaded to a private container registry.
3. A Pod configuration schema is defined/updated with details of the corresponding image's tag and any necessary Pod-specific settings/startup commands.
4. The Pod configuration schema is applied, either statically or dynamically. When the schema is applied, Kubernetes pulls the image from the container registry and creates a new Pod according to the schema, running any startup commands if provided.
5. If Pod creation is successful, the result is a new Pod running in the cluster.

1. Docker Container Creation Docker containers are created using a Dockerfile. The Dockerfile used to create the container image for the RR Message Server is shown below:

```
# Base Image
FROM ubuntu:latest
# Package Installation
WORKDIR /tmp/
```

```

ENV DEBIAN_FRONTEND="noninteractive"
RUN apt-get update && apt-get install -y \
python3-pexpect python3-pip
# User Creation
RUN useradd -ms /bin/bash rrserver
USER rrserver
# File structure creation/app setup
RUN pip3 install requests python-socketio \
eventlet
WORKDIR /home/rrserver/
RUN mkdir app
WORKDIR /home/rrserver/app/
COPY server.py .
COPY startup.sh .
# Startup command
CMD ["sh", "startup.sh"]

```

Listing 6: RR Message Server Dockerfile

The build process for each collaborative debugger Docker image follows the same structure as the one outlined in the Dockerfile above:

1. The base image is defined. The RR Message Server and RR Debug Session images are based on the latest Ubuntu image. This is particularly necessary for the RR Debug Session image, as rr's low-level nature necessitates somewhat frequent updates as changes are made to the Linux kernel. The Frontend/API Server image is based on the latest Node image, and the Database image is based on the latest MongoDB image.
2. Second, any necessary packages are installed. For the RR Debug Session image, rr is compiled from source and installed.
3. A non-root user is created if necessary.
4. Program files are copied over and a file structure is created. Pack-

ages that don't rely on the base image's built in package manager are installed at this time. In the example above, these include Python packages.

5. A startup command is defined.

Each line in a Dockerfile corresponds to a layer in the built image. This build order minimizes the amount of rebuilding necessary by placing the items that are most likely to change towards the end of the build process.

2. Container Registry Upload Most Managed Kubernetes services come with the option to create a private container registry. With proper authentication, this allows Docker and Kubernetes to access user-created images as easily as if they were in a public registry. Images built with Docker are uploaded to a private container registry for use in the collaborative debugger.

3. Pod Configuration Schema The Pod configuration schema for most Pods in the collaborative debugger is fairly generic. It consists of a `name`, an `image` sourced from the container registry, and in the case of Pods that need to interact with a load balancer, an `app`.

```
apiVersion: v1
kind: Pod
metadata:
  name: rr-translation
  labels:
    purpose: translate-rr
spec:
  containers:
  - name: rr-test-container
    image: example-container-registry.com/sproj/...
```

```
securityContext:
  capabilities:
    add:
      - SYS_PTRACE
restartPolicy: OnFailure
```

Listing 7: RR Debug Session Schema

A notable exception is the RR Debug Session Schema, which adds the `SYS_PTRACE` capability to the Pod. This is necessary for `rr` to properly trace system calls.

4 & 5. Pod Creation For statically created Pods, the `kubectl apply` command is used to create new Pods. Kubernetes pulls the container image defined in the schema from the container registry and starts the Pod with any necessary commands. The Database Pod is connected to a long-term storage volume at this time. Upon successful creation, the Pod is ready to interact with any necessary load balancers.

5.3.3 Service Creation

Services are the first components of the collaborative debugger created after the cluster is initialized. The three Services used by the debugger are all statically created. Like statically created Pods, schemas that define Services are applied manually using the `kubectl apply` command.

```
apiVersion: v1
kind: Service
metadata:
  name: frontend-load-balancer
spec:
  selector:
    app: frontend
```

```
type: LoadBalancer
ports:
  - port: 3000
    targetPort: 3000
```

Listing 8: Frontend Load Balancer Schema

The `app` field in the above schema corresponds to the `app` field defined in the metadata of the Frontend/API Server configuration schema. Traffic to the Load Balancer’s external IP address on port 3000 is redirected automatically by Kubernetes to any Pod running the Frontend/API Server. The Load Balancer determines which Pod is the most suitable given current demands on the system.

6 RR Debug Sessions & RR Message Server

At the heart of the collaborative debugger is the RR Debug Session. Every time a user wants to debug a new program execution, an RR Debug Session Pod is dynamically created by the API server. The new Pod is assigned a unique five digit identification number when it is created. This five digit number is used as the channel for the RR Debug Session, separating its communication from other RR Debug Session Pods. Clients, RR Message Server Pods, and the RR Debug Session Pod connect through the RR Message Load Balancer. Clients and RR Debug Sessions send messages using the channel assigned to the Pod at creation time. A diagram of one full communication cycle between two clients and an RR Debug Session Pod is outlined in the diagram below:

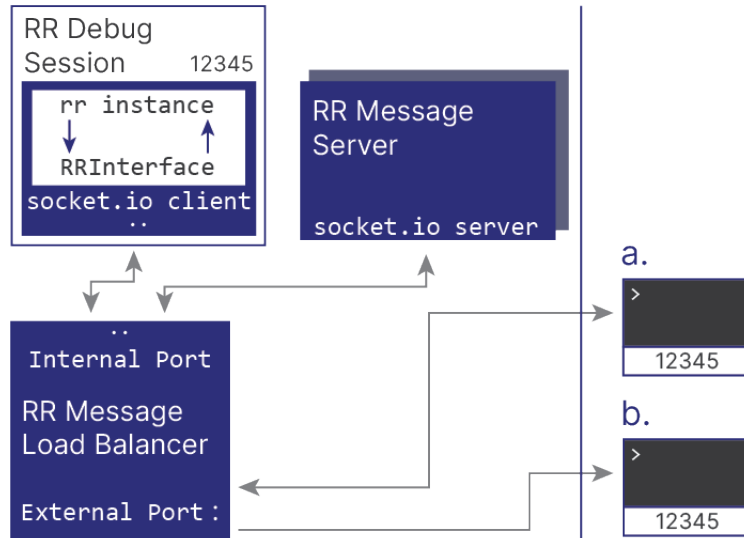


Figure 9: RR Message Server Communication Cycle

The process of sending a command to an RR Debug Session and receiving a response is as follows:

1. Client A, already having connected to the channel that corresponds to its current RR Debug Session (channel 12345 in the example), sends an rr command.
2. The RR Message Server receives the command and emits a message that RR Debug Session Pods are equipped to receive on the same channel. In practice, since only one RR Debug Session Pod is ever on a channel, this equates to emitting a message to the Pod directly.
3. The RR Debug Session Pod receives the message, and passes the command to its instance of RRInterface, the class which controls rr. When it receives a response, it emits the output from rr along with other debugging information and the channel.
4. The RR Message Server receives the response, and emits a response message that clients are equipped to receive on the corresponding channel.
5. Clients A and B are connected to the corresponding channel, so they both receive the response.

6.1 RR Message Server

6.1.1 Tools Used

The collaborative debugger needs an efficient way for clients to communicate with running RR Debug Sessions. The Socket.IO library was selected to construct this efficient communication pathway.

Socket.IO To speed communication, the collaborative debugger uses WebSockets to directly connect web clients and the Pods running rr. Socket.IO is a library that extends WebSockets. It provides backup in case a WebSocket connection cannot be established, enables automatic reconnection and disconnection detection, and adds support for namespaces [16]. The collaborative debugger uses the standard JavaScript implementation of Socket.IO on the client side. Messages are passed through a server to individual debugging Pods, both of which use the Python implementation of Socket.IO, `python-socket.io` [11].

6.1.2 Overview

The RR Message server is remarkably simple, consisting of just 3 important functions:

```
@sio.on('join_channel')
def join_channel(sid, data):
    sio.enter_room(sid, data['channel'])
@sio.on('rr_command')
```



```

def on_rr_command(sid, data):
    data['sid'] = sid
    try:
        sio.emit('rr_command', data,
                  room=data['channel'])
    except:
        pass

@sio.on('rr_response')
def on_rr_response(sid, data):
    sio.emit('rr_response', data,
              room=data['channel'])

```

Listing 9: RR Message Server

Even in a fully production ready version of the collaborative debugger with more security features enabled, the RR Message Server is unlikely to become much more complex. It exists purely to pass messages between clients and RR Debug Session Pods and to manage channels. Since the design of other aspects of the collaborative debugger ensure that all members of a given channel exist only during the lifetime of it's corresponding RR Debug Session Pod, the `'join_channel'` event handler simply adds socket.io clients to a specified channel on request.

The two message processing functions, `on_rr_command` and `on_rr_response`, are equally simple. When the RR Message Server receives an `'rr_command'` message, it passes the message data along to the corresponding RR Debug Session Pod by emitting an `'rr_command'` namespaced to the correct channel with `room=data['channel']`. Since the socket.io client running in RR Message Server Pods has an event handler defined for `'rr_command'`'s but frontend clients do not, only RR Message Server Pods receive the command. The same process happens in reverse for `'rr_response'`'s, with only the

frontend socket.io clients having a handler defined for `'rr_response'`.

6.2 RR Debug Session

6.2.1 Tools Used

RR Debug Sessions primarily manage and communicate with rr, which has already been covered in detail in 4.1.1. pygdbmi is used for this process.

pygdbmi In order to “support the development of systems which use the debugger as just one small component of a larger system”, gdb provides a machine-oriented interface called gdb/mi [22]. rr supports interaction through gdb/mi, and using the interface was a natural choice for the collaborative debugger. In addition to being far easier to interact with from within a program, the structured, machine-friendly output of gdb/mi lends itself in particular to future development of visualization aids in the collaborative debugger.

To parse rr output into Python dictionaries and to easily control rr as a subprocess, pygdbmi [10] is used in each debugging Pod. pygdbmi’s abstraction simplifies programmatically controlling rr. A Pod can receive a command from the client, pass it to rr, and respond without having handle directly parsing gdb/mi output or managing the rr process.

6.2.2 Building Example Pods

The process for building RR Debug Session example images is differs slightly from other images used by the collaborative debugger. Currently, three example RR Debug Session images are available for use: **hash**, **smash**, and **memoization**. Example container images are based on the primary RR Debug Session image, RR Translation. The build process for this image, based on the lasted official Ubuntu container image, installs rr, python, and all necessary packages as well as the app that will run on the final image, `rrtranslation.py`.

```
FROM example-container-registry.com/sproj/...
WORKDIR /home/debug/app/
COPY hash.c .
COPY names .
COPY startup.sh .
```

Listing 10: RR Debug Session Hash Example—Dockerfile

The Dockerfile shown above is for the **hash** RR Debug Session example. The build process copies over any necessary files, as well as a startup script:

```
gcc -g -o hash hash.c
rr record ./hash
python3 rrtranslation.py $1
```

Listing 11: Example Startup Script

This script is used to start the '`rrtranslation.py`' program when the Pod is dynamically created. The script will be passed the Pod's channel as argument `$1` by the Kubernetes API on startup.

Current example Pod startup scripts compile the program to be debugged

and record execution when the Pod is created.

6.2.3 The RR Translation Program

The RR Translation program consists of two main components: a Socket.IO client, `sio`, and an instance of the `RRInterface` class, `rri`, that controls the rr subprocess and parses interactions. Simple code to read information about the current source file currently exists within `sio`'s `'rr_command'` event handler, but should be broken out into its own class if more complexity is added.

The RR Translation program's socket.io client instance begins by connecting to the RR Message Server. It immediately emits a `'join_channel'` message, using the channel number passed in by the Kubernetes API on Pod creation (7.4.3).

```
if __name__ == '__main__':
    channel = sys.argv[1]
    sio.connect('http://rr-message-server-load-balancer:8000')
    sio.emit('join_channel', {'channel': channel})
```

Listing 12: RR Translation Main

After joining the appropriate channel, `sio` waits to receive an `'rr_command'`. The first part of `sio`'s `'rr_command'` event handler is shown below:

```
@sio.on('rr_command')
def on_rr_command(data):
    body = {'from': data['sid'],
            'command': data['command'],
            'channel': channel}
    try:
```

```
response = {'output': rri.write(data['command'])}
```

Listing 13: RR Command Event Handler

The handler first builds part of the response body, passing back data about the channel and originator of the message that will be important for the RR Message Server and client later. Next, it calls the function necessary to pass a message to rr and receive a response, `rri.write()`.

When the program starts, an instance of the `RRInterface` class, `rri`, is initialized in the same scope as `sio`. `RRInterface` contains an instance of the `gdbController` class from `pygdbmi` to interface with rr and control the rr subprocess, as well as a collection of methods to parse rr output, the three most important of which are shown below:

```
def get_full_rr_response(self, command):
    response = self.gdbmi.write(command)
    while(not(self.end(response) or self.exited(response))):
        response.extend(self.get_rr_response())
    return response

def write(self, command):
    if self.command_forbidden(command):
        raise DissallowedError
    self.timeline.append(command)
    return self.console_output(
        self.get_full_rr_response(command))

def source(self):
    f = self.get_full_rr_response(
        '-file-list-exec-source-file')[0]['payload']
    return {'file': f['file'],
            'line': f['line'],
            'path': f['fullname']}
```

Listing 14: RRInterface

When `write` is called, it determines if the command to be executed is forbidden. Currently only one gdb command, `shell`, is disallowed. Parsing

the output from `shell` is difficult, it is of virtually no use since *recordings*, not currently executing programs that `shell` could effect, are being debugged. It also opens the door to security issues. Though containerization means that the most users could probably do is render their own debug session useless, the downsides of `shell` far outweigh the benefits.

Next, `write` appends the command to the `RRInterface`'s `timeline` instance variable. `timeline` is a list of all commands the debug session has executed. Though unused at the moment, it can serve in the future to implement a shared history between all clients and to facilitate the saving of debug sessions in-progress. By issuing all commands in `timeline` to a new instance of `RRInterface` with the same recording, it is possible to restore a debug session to an identical previous state. To save a debug session, the recording and `timeline` can be stored in a database, and be used to initialize a new Pod with identical state to a previous RR Debug Session Pod. This appears to clients as a seamless restoration of a previous Debug Session.

Since the frontend currently consists of a terminal interface to rr and a view of the current source code, `write` does not need to return any information from rr other than console output. Future versions of the collaborative debugger that support visualization tools could use gdb/mi commands such as `-stack-list-frames` to retrieve information to be used by frontend visualization tools. `console_output` extracts user-relevant output from the dictionary returned by `GDBController`'s `write` method. The `get_full_rr_response` method ensures that all relevant output has been

received from the `rr` subprocess before returning. `write` reduces the lengthy dictionary that would be produced by a command as simple as `break main` into a few lines of user-relevant output.

After the `'rr_command'` event handler has received a response from `rri`, it calls `RRInterface`'s `source` function to retrieve information about the current source file being debugged. Most often the source file has not changed, and the only relevant piece of information that needs to be passed back to the client is the new line number in the source file. If the source file has changed, the handler reads its contents, updates `current_file`, and emits the line number, file name, and contents. If an error occurs at any point in the process, a detailed trace is emitted for debugging purposes. The trace should be omitted in production.

6.3 Considerations

Most of the complexity in the RR Translation program stems from parsing `rr` output. The flow of data in the program is quite simple: a command is received, it is passed to `rr`, `rr`'s response is processed, and a response message is emitted. This flow would remain unchanged even with the addition of functions to retrieve information for data visualization. Since the process of turning `rr` responses into a terminal interface takes place in the frontend, rather than the RR Translation program providing some sort of REPL over WebSockets, updates can be made to the frontend webapp without requiring backend changes. The whole process is also quite snappy, with only a slight

delay from the client's perspective compared to running a debugger natively.

Effort is taken in the design of the RR Debug Session and RR Message Server to ensure students receive a pleasant and near-identical experience to debugging locally in the future. The system of inheritance used to create example Debug Session Pods is designed to be as simple as possible, though a more fully featured frontend could streamline this process further in the future.

7 Frontend/API Server

The second major backend system consists of the Frontend/API Server and database. This system is responsible for serving the frontend webapp, as well as processing all API requests, such as those to authenticate users and create new RR Debug Sessions.

7.1 Tools Used

7.1.1 MongoDB

The collaborative debugger uses a database to store information about users, Pods, and example debugging sessions. Due to its speed of deployment and natural interaction with the object-oriented languages used to create the project, MongoDB was chosen as database software [7].

7.1.2 Flask and Node.js

The primary server for the collaborative debugger is split into two sections: a simple Node.js [9] server that serves the frontend webapp, and an API server created using Flask [2]. While in development, the builtin React (4.1.2) development server is used to serve the frontend. This makes debugging the frontend far easier.

An API server is necessary to authenticate users and to provide a means to create/delete debugging sessions. Since the rest of the backend was cre-

ated using Python, Flask was chosen to create the API server. Flask is a lightweight web application framework which lends itself perfectly to interacting with the Python MongoDB and Kubernetes APIs.

7.2 Overview

The builtin React development server is used to serve the frontend webapp and redirect API requests to the Flask API server during development. For both security and performance reasons this should be changed to a combination of a more robust solution like Nginx [8] and a production suitable web server for production. The relationships between components and overall process will remain unchanged after this migration.

Since the entire webapp is one page, the process for serving it is completely standard. The server receives a request for the index page of the website, and returns the compiled React webapp that makes up the homepage. The process for API requests is more involved, and an example is outlined below:

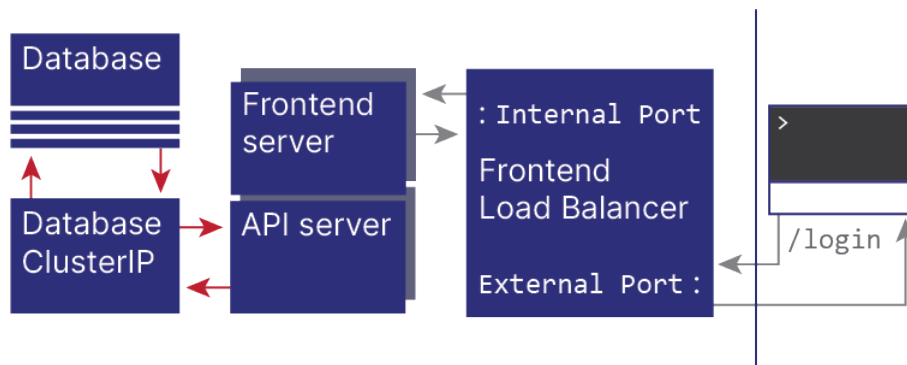


Figure 10: API Request Process—Login

1. The client makes a POST request to the `'/login'` URL. This request contains the necessary data to log in the user.
2. The frontend server forwards the request of a URL it does not recognize to the API server running in the same Pod on a different port.
3. The API server recognizes the `'/login'` route, and communicates with the database to attempt to log in the user.
4. The API server returns the result of the login request, which is returned by the frontend server to the client.

7.3 Frontend/API Server Configuration

The Frontend/API Server uses Yarn [18] to manage packages. When a Frontend/API Server Pod is deployed, the Pod's startup script uses the `yarn start` and `yarn start-api` commands to start the frontend server and API server respectively. These scripts are defined in the server's `package.json` configuration file:

```
"start": "react-scripts start",
"start-api": "cd api && flask run --no-debugger",
.
.
.
"proxy": "http://localhost:5000",
```

Listing 15: Frontend/API Server Configuration

Also of note in `package.json` is the instruction to proxy the API server, which is running on port 5000. This achieves the automatic forwarding of

API requests to the API server described above. Since API requests are proxied through the same address serving the frontend webapp, there are no issues with cross-origin requests.

7.4 The API Server

The API server is implemented using Flask. The server program consists of handlers for the various API routes and instances of two classes which communicate with the database and/or Kubernetes API: `TranslationPodManager` and `UserManager`. The full structure of the API Server program, `api.py` is shown below:

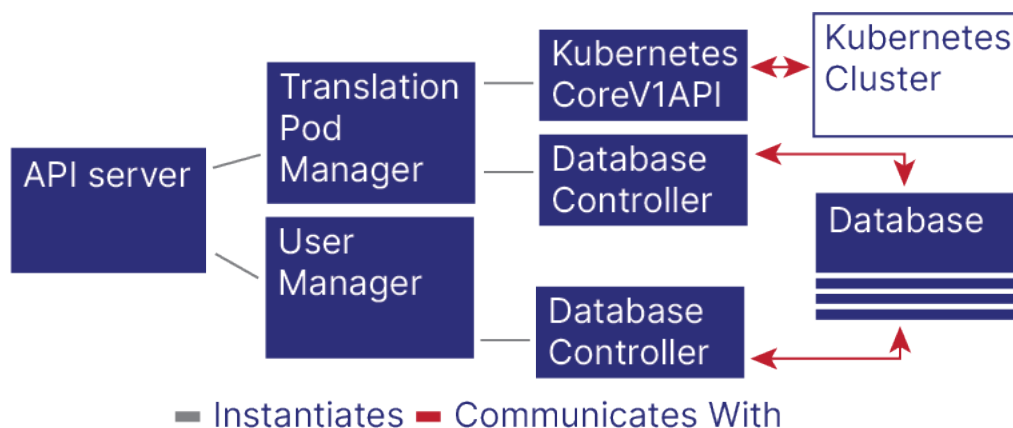


Figure 11: Structure of the API Server

Care was taken to abstract as much as possible in the design of the API server and the classes it instantiates. The two classes instantiated by `api.py`, `TranslationPodManager` and `UserManager`, each instantiate another class which interacts with the database, `DatabaseController`, as an instance vari-

able. `TranslationPodManager` further instantiates the Kubernetes Core V1 Python API to communicate with the cluster in order to create/delete Pods. This process of abstraction ensures that the database, underlying cluster structure, and authentication methods can all change without significant changes needing to be made to `api.py`.

The most significant improvement that could be made to this structure would be to break the `UserManager` and `TranslationPodManager` instances out into separate programs running on their own Pods in the cluster. This would allow changes to be made in those classes, say to update the process of deleting Pods, without needing to update the entire Frontend/API Server Pod. For the time being, the relative simplicity of the API means that the added complexity and overhead of implementing a method to communicate between dedicated `TranslationPodManager` and `UserManager` Pods and Pods running the Frontend/API Server doesn't seem worthwhile. This may change in the future.

Currently supported API requests are as follows:

1. `/login`: Logs in the user given by `'name'`, or, (given the rather lax development security) creates a user if none matching `'name'` exists. Returns the name of the user.
2. `/channel`: How users join RR Debug Sessions. Binds the user given by `'name'` to an RR Debug Session Pod matching `'channel'` in the database, if such a Pod exists. Returns the channel.

3. `'/pods'`: How the client gets a list of RR Debug Session Pods the user is currently participating in. Returns the channels for all RR Debug Session Pods who the user, given by `'name'`, is currently bound to in the database.
4. `'/examples'`: How the client gets a list of example RR Debug Session Pod names. Returns all names of example RR Debug Session Pods that exist in the database.
5. `'/new'`: Creates a new RR Debug Session Pod based on the name given by `'program'`. Binds the user, given by `'name'`, to the Pod. Returns the channel of the Pod. This request is the most complex, and is covered in more detail below.
6. `'/delete'`: Deletes a binding between the user given by `'name'` and the RR Debug Session Pod given by `'channel'`. If there are no bindings left between the Pod and users (if all the users have left the debug session), deletes the Pod from the database and the cluster. Returns `True`.

All the API request handlers listed above return an error message in the case of an error occurring during the processing of an API request. Since most errors that could occur would either be the result of unanticipated issues on the part of the user, or of some unforeseen catastrophic internal error, informing users of error specifics is often unhelpful. Some `TranslationPodManager`

and `UserManager` functions throw specific errors in the event of duplicate usernames, channels that do not exist, etc. These errors are caught and meaningful error messages are returned to the frontend webapp, which can then pass them on to the user.

To examine the process of processing an API request, it makes sense to look at the most complex example, creating a new RR Debug Session Pod. This example shows the process for processing an API POST request, interacting with the database, and dynamically creating a Pod. All API requests follow a similar procedure.

7.4.1 Processing a POST Request—Extracting Information

Below is the route handler for the `’/new’` URL, as well as the instantiation of `TranslationPodManager` and `UserManager`.

```
tpm = podmanager.TranslationPodManager(
    url='mongodb://database-load-balancer',
    port=27017)
um = usermanager.UserManager(
    url='mongodb://database-load-balancer',
    port=27017)
app = Flask(__name__)
@app.route('/new', methods=['POST'])
def new():
    try:
        name = request.get_json()['name']
        program = request.get_json()['program']
        channel = tpm.create_pod([name], program)
        return {'channel': channel}
    except:
        return {'error': 'Internal Error'}
```

Listing 16: API Server New RR Debug Session Event Handler

All API routes only support POST requests. The route handler first extracts the relevant information from the POST request body, in this case `name` and `program`. `name` always corresponds the user who is making the request's username, and `program` corresponds to the name of the example RR Debug Session image to be used.

`new` then calls `TranslationPodManager`'s `create_pod` function to create a new RR Debug Session Pod. `create_pod` can bind multiple users to a Pod when it is created, so `name` is passed inside a list.

7.4.2 The Database

Before drilling into `create_pod`, it may be helpful to examine the MongoDB database which `TranslationPodManager` and `UserManager` interact with. The database consists of three collections: `users`, `pods`, and `examples`. Below is an example record from each collection:

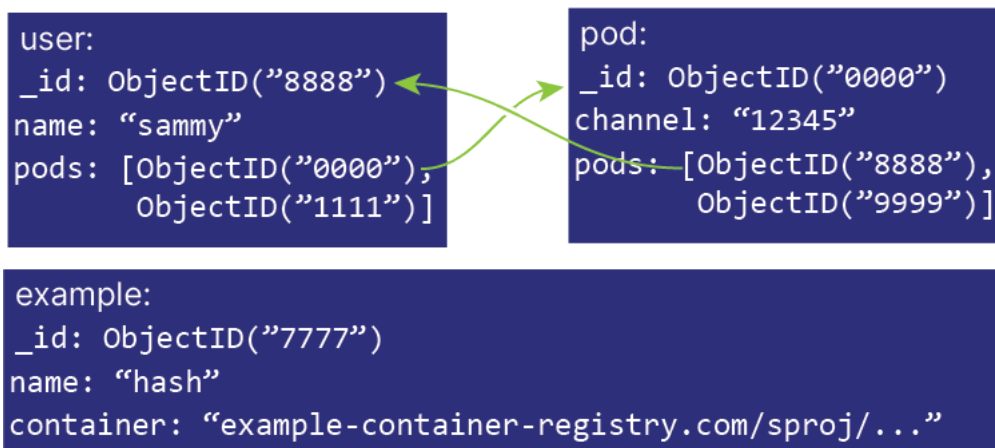


Figure 12: Database Record Examples

`Pods` and `Users` have a many-to-many relationship. Each user can be bound to an unlimited number of pods, and vice versa. The `DatabaseController` class includes functions to efficiently retrieve `Users` and `Pods` given one another, and to create/destroy bindings between `Users` and `Pods`.

`Examples` are independent of `Users` and `Pods`. The container registry URLs of RR Debug Session example images are stored in the database so that new images can be easily added to the collaborative debugger upon creation.

MongoDB keeps `_ids` unique. In addition, the collaborative debugger defines uniqueness constraints on `Users.name`, `Pods.channel`, and `Examples.name` when the database is created.

7.4.3 Processing a POST Request—Dynamic Pod Creation

To dynamically create a Pod, `create_pod` gets the container registry URL of the image to base the Pod on, creates the new Pod, binds users to the Pod, and returns the channel the Pod was created with. The first part of the function is shown below:

```
def create_pod(self, names, program):
    examples = self.get_examples()
    image = None
    try:
        image = examples[program]
    except:
        raise NoSuchExampleError
```

Listing 17: Pod Creation 1

`create_pod` calls `TranslationPodManager`'s `get_examples` method to

get a list of RR Debug Session example image names and container registry URLs from the database. In the event that the user has somehow passed a spurious image name or has passed a name when there are no images, an exception is raised. Otherwise, a RR Debug Session Pod schema is created using the image:

```
dep = {'apiVersion': 'v1',
      'kind': 'Pod',
      'metadata': {'labels':
                    {'purpose': 'translate-rr'}}},
      'spec': {'containers':
                [{ 'image': image,
                  'name': 'rr-test-container',
                  'command': ['sh'],
                  'args': ['startup.sh'],
                  'securityContext':
                    {'capabilities':
                     {'add': ['SYS_PTRACE']}}}]},
                'restartPolicy': 'OnFailure'}}
```

Listing 18: Pod Creation 2

This schema, represented as a dictionary in Python, is identical to the YAML representation of the RR Debug Session Pod creation schema shown in (5.3.2) with the exception of `args` which will be updated further in a later step. Finally, `create_pod` does the work necessary to dynamically create a Pod:

```
channel = self.dbc.add_pod(
    self.dbc.get_userids_by_name(names))
dep['metadata']['name'] = channel
dep['spec']['containers'][0]['args'].append(channel)
resp = self.v1.create_namespaced_pod(
    body=dep, namespace='default')
# Wait to return the channel until the pod is live
# and read to receive incoming communication
status = self.v1.read_namespaced_pod_status(
    channel, 'default').status.container_statuses
```

```
while status == None or status[0].state.running == None:
    time.sleep(1)
    status = self.v1.read_namespaced_pod_status(
        channel, 'default').status.container_statuses
return channel
```

Listing 19: Pod Creation 3

First, `DatabaseController`'s `add_pod` method is called to insert a new `Pods` record into the database. `add_pod` randomly generates a new five digit `channel` for the Pod, using the uniqueness constraint imposed on `Pods.channel` to ensure an unused `channel` is generated. If no channels are open an error is raised. Otherwise, the users passed to `add_pod` are bound to the new Pod in the database, and `channel` is returned. This process has the effect of limiting the number of simultaneous RR Debug Sessions that can be in use at the same time to 89,999, and of slowing as more channels are used. Given the few current users, the likelihood of a collision when generating a new channel is low enough that a more advanced method is unnecessary.

Next, the name of the container is updated to `channel`. In addition, the startup arguments are updated to include `channel`. This is how `channel` is passed to the `startup.sh` script and eventually used to join a channel on the RR Message Server in `rrtranslation.py` (6.2.2 & 6.2.3).

The Kubernetes API is then finally used to create the Pod. The Pod is created in the default namespace. This could be changed to logically isolate dynamically and statically created Pods with minimal hassle, since Kubernetes' DNS resolution can connect Pods across namespaces. If Pod creation fails, an error will be thrown. Otherwise, a loop is used to wait

until `rrtranslation.py` is running inside the Pod. Once the Pod is ready, `channel` is returned.

7.4.4 Processing a POST Request—Returning a Response

Most request handlers, `’/new’` included, simply return a JSON serializable dictionary containing whatever results the API Server’s instance of `TranslationPodManager` or `UserManager` returned, or an error message. The `’/new’` handler returns `’channel’: channel` so that the frontend webapp’s `socket.io` client can join the channel corresponding to the RR Debug Session Pod that was just created.

7.5 Considerations

The Frontend/API server connects all components of the collaborative debugger together. It authenticates users through the database, joins the frontend and Kubernetes cluster through the dynamic creation of Pods, and starts the interaction between users and the programs they wish to debug. Care has been taken to make it modular and extensible within reason, and additional authentication or Pod creation features should be easy to add.

The experience provided by invisible components such as the Frontend/API Server is just as important as that provided by the frontend webapp. All components of the collaborative debugger hopefully come together to provide an experience that facilitates the teaching and learning of debugging.

8 Next Steps

8.1 Features to Add

The current iteration of the collaborative debugger successfully enables collaboration between students. Significant effort has been taken to design the debugger in a way that makes collaboration seamless. The collaborative debugger hopefully fulfills it's first requirement—to encourage collaboration.

Though the collaborative debugger makes it easy for teachers to design and distribute debugging lessons, more work is necessary to fully fulfill the second goal—that of enhancing students' understanding of debugging and of programs being debugged.

The most significant addition towards this goal would be that of frontend visualization tools. A simple visualization of the stack and registers would greatly expand students' understanding of the program being debugged in all three example sessions currently included with the collaborative debugger. `gdb/mi` supports commands which give structured output of both these aspects of the execution space, and the frontend has been designed to make adding visualizations of them simple. Adding visualization features will bring the collaborative debugger much closer towards hopefully fulfilling its second requirement.

8.2 Testing Efficacy

Though a review of the research seems to indicate that a tool like the collaborative debugger should help make students more effective at debugging, further study is needed to draw conclusions. First, more research is necessary to determine definitively if teaching debugging collaboratively is more effective than teaching it individually. Second, the collaborative debugger needs to be thoroughly tested to determine whether it is an ideal platform for collaborative debugging.

A study similar to those undertaken in [23], [26], and [27] could be done to test both of these aspects. Students participating in an introductory systems programming class could be split into three groups. Group one could study debugging individually, using `rr` by themselves to progress through debugging sessions designed by the instructor. Group two could study debugging in pairs, relying on existing tools such as video-conferencing software to facilitate collaboration. In the final group, pairs would use the collaborative debugger to work together on the same debugging sessions as groups one and two. Students' individual and pair performance in debugging both unseen and self-generated programs could then be tested.

8.3 Conclusions

The collaborative debugger aims to encourage collaboration and build students' knowledge of debugging through careful design. Many components

of a distributed system come together to create a platform that hopefully benefits both teachers and students. Through deliberate design, the collaborative debugger can conceivably be the ideal platform to teach and learn debugging, both remotely and in-person.

References

- [1] Docker. <https://www.docker.com/>.
- [2] Flask. <https://github.com/pallets/flask>.
- [3] Kubernetes. <https://github.com/kubernetes/kubernetes>.
- [4] Kubernetes Documentation. <https://kubernetes.io/docs/home/>.
- [5] Managed Kubernetes on DigitalOcean.
<https://www.digitalocean.com/products/kubernetes/>.
- [6] Monaco editor. <https://github.com/microsoft/monaco-editor>.
- [7] MongoDB. <https://www.mongodb.com>.
- [8] Nginx. <https://nginx.org/en/>.
- [9] Nodejs. <https://github.com/nodejs/node>.
- [10] pygdbmi. <https://github.com/cs01/pygdbmi>.
- [11] python-socketio. <https://github.com/miguelgrinberg/python-socketio>.
- [12] React. <https://github.com/facebook/react>.
- [13] react-monaco-editor. <https://github.com/react-monaco-editor/react-monaco-editor>.
- [14] react-xterm. <https://github.com/farfromrefug/react-xterm>.

- [15] Replit. <https://repl.it/>.
- [16] socket.io. <https://github.com/socketio/socket.io>.
- [17] Xterm.js. <https://github.com/xtermjs/xterm.js>.
- [18] Yarn. <https://yarnpkg.com/en/>.
- [19] BROWN, N. C. C., SENTANCE, S., CRICK, T., AND HUMPHREYS, S. Restart: The resurgence of computer science in uk schools. *ACM Trans. Comput. Educ.* 14, 2 (June 2014).
- [20] CHMIEL, R., AND LOUI, M. C. Debugging: From novice to expert. In *Proceedings of the 35th SIGCSE Technical Symposium on Computer Science Education* (New York, NY, USA, 2004), SIGCSE '04, Association for Computing Machinery, p. 17–21.
- [21] DOLAN-GAVITT, B., HODOSH, J., HULIN, P., LEEK, T., AND WHELAN, R. Repeatable reverse engineering with panda. In *Proceedings of the 5th Program Protection and Reverse Engineering Workshop* (New York, NY, USA, 2015), PPREW-5, Association for Computing Machinery.
- [22] FREE SOFTWARE FOUNDATION. *Debugging with GDB*, 2020.
- [23] HANKS, B., MCDOWELL, C., DRAPER, D., AND KRNJAJIC, M. Program quality with pair programming in cs1. *SIGCSE Bull.* 36, 3 (June 2004), 176–180.

- [24] LI, C., CHAN, E., DENNY, P., LUXTON-REILLY, A., AND TEMPERO, E. Towards a framework for teaching debugging. In *Proceedings of the Twenty-First Australasian Computing Education Conference* (New York, NY, USA, 2019), ACE '19, Association for Computing Machinery, p. 79–86.
- [25] MCCAULEY, R., FITZGERALD, S., LEWANDOWSKI, G., MURPHY, L., SIMON, B., THOMAS, L., AND ZANDER, C. Debugging: a review of the literature from an educational perspective. *Computer Science Education* 18, 2 (2008), 67–92.
- [26] MCDOWELL, C., WERNER, L., BULLOCK, H. E., AND FERNALD, J. Pair programming improves student retention, confidence, and program quality. *Commun. ACM* 49, 8 (Aug. 2006), 90–95.
- [27] MICHAELI, T., AND ROMEIKE, R. Improving debugging skills in the classroom: The effects of teaching a systematic debugging process. In *Proceedings of the 14th Workshop in Primary and Secondary Computing Education* (New York, NY, USA, 2019), WiPSCE'19, Association for Computing Machinery.
- [28] O'CALLAHAN, R., JONES, C., FROYD, N., HUEY, K., NOLL, A., AND PARTUSH, N. Engineering record and replay for deployability: Extended technical report. *CoRR abs/1705.05937* (2017).

- [29] O'CALLAHAN, R., JONES, C., FROYD, N., HUEY, K., NOLL, A.,
AND PARTUSH, N. rr, Dec. 2019. <https://github.com/mozilla/rr>.